

# As If, Not As Was

## MARS-Bench and the Agentic Audition of LLMs for Marcionite Surrogation

A.G. Elrod

Vrije Universiteit Amsterdam, PhD candidate

HZ University of Applied Sciences, docent-onderzoeker

May 12, 2026

### Abstract:

Marcion of Sinope has been silent for eighteen centuries, and what survives of his voice has been preserved primarily through the writings of his ancient opponents. Any project that would reason in his idiom faces a constitutive absence: the historical Marcion *as he was* cannot be recovered. This paper proposes that the absence can nonetheless be operationally staged. Drawing on Stanislavski's distinction between actuality and *scenic truth* and on Joseph Roach's account of *surrogation* (the cultural process by which communities fill vacancies left by death or departure) it reframes the selection of a Large Language Model for historical-theological work as an *audition*: not a test of resemblance, but a test of eligibility to perform within disciplined given circumstances. MARS-Bench is the instrument that performs this audition. Spanning seven domains (ancient languages and philology, biblical knowledge, patristic knowledge, early Christian history, theological knowledge, heresiological knowledge, and Marcion Studies) and governed through frozen reference answers, approved context packets, response-integrity checks, and paired agentic judging using Codex and Claude Code, the benchmark evaluates fourteen candidate models across 251 questions and 10,542 responses. Inter-judge reliability is high ( $ICC(A,2) = 0.962$ ); the role-weighted composite, supplemented by run-to-run variance, license, and architecture diagnostics, supports gemma-4-31b-it as the provisional substrate for downstream Synthetic Theological Agent work, with llama-3.3-70b-instruct retained as comparator. MARS-Bench validates a selection procedure, not a resurrected voice. Marcion *as was* remains beyond reach; Marcion *as if* is operationally stageable, and that distinction is the methodological contribution of the audition.

# Introduction

*Totus mundus agit histrionem.*  
The motto of Shakespeare's Globe Theatre

Russian actor, director, and theorist Konstantin Stanislavski fundamentally shaped modern performance with a discipline whose purpose was not theatrical mimicry but truthful inhabitation. In *An Actor Prepares*, he presents the actor's work as a conscious and unconscious technique for preparing the conditions under which one may *live the part* afresh, not by forcing inspiration, but by moving through action, attention, imagination, and the role's "given circumstances."<sup>1</sup> Those circumstances, the story, the social world, the authorial text, the directorial interpretation, the actor's own experience, do not merely restrict the actor. They liberate the actor by giving the actor something to *live*. The actor's task is, therefore, neither costume nor ventriloquism, neither raw self-expression nor antiquarian reproduction, but a disciplined act of substitution: a living body and consciousness enter a space prepared by words, memory, absence, and expectation, making present what is not present in a literal sense.

That is the first sense in which this paper means *as if, not as was*. The "if" is not a mere pretense, but a rule-bound imaginative condition. It authorizes action inside a world one did not make. Joseph Roach's later concept of *surrogation* names the broader cultural process by which communities attempt to fill vacancies left by death, departure, or loss; Stanislavski's system gives a craft vocabulary for how such substitution can become scenic truth rather than mere costume or caricature.<sup>2</sup> The present paper transposes that performance-theoretical grammar into a setting that may initially seem strange: the agentic benchmarking of large language models (hereafter LLMs) for historical-theological work. Its question is not whether an LLM can resurrect Marcion of Sinope, recover his mind, or reproduce the second century "as was." It asks, more modestly and more technically, how one might audition models

---

<sup>1</sup> Constantin Stanislavski, *An Actor Prepares*, trans. Elizabeth Reynolds Hapgood (New York: Routledge, 1989; orig. 1936), 15–17, 54–56, 302–3. See also the active draft's note that "if" is a significant concept in Stanislavski's system: for an actor inhabiting a role within given circumstances, "if" gives the actor a motive for action.

<sup>2</sup> Joseph Roach, *Cities of the Dead: Circum-Atlantic Performance* (New York: Columbia University Press, 1996), esp. 2–3, 21, 36. Roach's account is historically situated in circum-Atlantic performance culture rather than offered as a universal transhistorical mechanism; this paper therefore transposes rather than simply imports his concept of surrogation.

for their capacity to answer *as if* within disciplined, source-critical, historically mediated given circumstances.

A casting story for the Ray Charles biopic *Ray* (2004) is instructive here because it gives the method its scene. When Jamie Foxx auditioned for the leading role, director Taylor Hackford brought him into Charles's own recording studio. Two electric pianos were set side by side. Charles sat at one, Foxx at the other. Charles first moved through funk and blues standards, which Foxx, who had played piano from childhood and later studied on a classical-piano scholarship at United States International University in San Diego, handled easily.<sup>3</sup> Then Charles shifted into Thelonious Monk. Foxx fumbled for what he later remembered as an unbearable stretch, with Charles pushing him, refusing to let the test become comfortable.<sup>4</sup> Hackford's interpretation of the scene is uncannily relevant to the argument of this paper. He observed, "Ray's not just testing him as a piano player, he's testing him as a man. And finally Jamie gets it. Ray stands up and hugs himself — he had a way of doing that when he got excited — and said, 'The kid's got it!'" Hackford's gloss followed: "I had cast Jamie as Ray Charles, but it was at that moment Ray anointed him."<sup>5</sup>

At the heart of the biopic's ability to transform historical actuality into dramatic truth is embodiment. Dennis Bingham observes that *Ray* depended upon an improbable convergence of musical and acting talent, concluding more generally that Hollywood biopics "do demand embodiment."<sup>6</sup> Foxx's audition illustrates why. First, education is an eligibility criterion: no amount of costuming could have substituted for the musical formation that made Foxx auditionable. Second, the test occurred inside Charles's idiom rather than in a neutral acting space. Third, what Charles tested was not description but performance. A writer might characterize what Ray Charles would have played; only an actor could sit beside him and play. The distinction matters because biographical performance fails when it collapses into surface impersonation. Foxx named this directly in an NPR interview: "Impersonation will kill you in a biopic." The work, he said, was to move past impersonation until the character was "locked down" enough that one could "improv within that."<sup>7</sup>

---

<sup>3</sup> Lola Ogunnaike, "Jamie Foxx Goes Dark to Play a Musical Hero," *New York Times*, 12 September 2004, Section 2, p. 33; Gloria Goodale, "An Audition for 'Ray' with Ray," *Christian Science Monitor*, 29 October 2004.

<sup>4</sup> *Ibid.*, see also Goodale, "An Audition for 'Ray' with Ray."

<sup>5</sup> Ogunnaike, "Jamie Foxx Goes Dark."

<sup>6</sup> Dennis Bingham, *Whose Lives Are They Anyway? The Biopic as Contemporary Film Genre* (New Brunswick, NJ: Rutgers University Press, 2010), 137–38.

<sup>7</sup> "Playing a 'Genius': Jamie Foxx Takes on Ray Charles," *NPR*, 22 October 2004. In the transcript, Foxx's negation of impersonation and his affirmative account of improvisation occur in one continuous

But what hath Hollywood to do with Athens—or, more precisely, with Pontus, Rome, and the hostile libraries of the early church? The analogy turns on an inversion. Foxx could be auditioned beside Ray Charles because Charles was still there. Every scholar who studies Marcion of Sinope faces a different and more difficult stage: unlike *Ray*, we cannot bring Marcion into the studio. He has been silent for eighteen centuries, and what survives of his voice has been preserved primarily through the writings of his ancient opponents, filtered through heresiological, polemical, and reconstructive traditions rather than through an independent Marcionite archive. Secondary scholarship has done impressive work in theorizing plausible Marcionite positions, but the field remains a contested reconstruction space rather than a stable script waiting to be performed.<sup>8</sup>

This is the problem to which my broader PhD project, “Interpreting the Unspoken: Counterfactual Historical Reasoning for Reconstructing Early Christian Heterodox Thought Using Large Language Models,” responds. The project asks whether LLMs, trained and constrained by carefully curated primary and secondary corpora, can be developed into Synthetic Theological Agents (hereafter, STAs): not recovered Marcionites, and not machines possessed by the past, but synthetic actors capable of surrogating Marcionite theological reasoning within specified historical, philological, and theological circumstances.<sup>9</sup> The STAs are the later performance. MARS-Bench is the audition.

That audition is necessary because base-model selection is not a neutral technical preliminary. A model that can recite summaries of Marcion from familiar scholarship may still fail as a theological surrogate if it cannot reason across ancient languages, biblical interpretation, patristic polemic, early Christian history, heresiology, and the variant reconstructions of Marcion Studies. Conversely, a model with impressive general ability may be unusable if it hallucinates sources, collapses hostile testimony into fact, ignores domain-specific uncertainty, or performs a generic “heretic” rather than a historically constrained Marcionite surrogate. MARS-Bench, therefore, asks which candidate models are eligible for later STA work by testing whether they can

---

response; this draft preserves the transcript’s conceptual linkage rather than treating them as separate claims.

<sup>8</sup> On the hostile and mediated source problem, see Tertullian, *Adversus Marcionem* 4–5; Epiphanius, *Panarion* 42; Adamantius, *Dialogue*; and the variant-reconstruction framework represented by Harnack 1921/1990, Hoffmann 1984, Schmid 1995, Kinzig 2004, Moll 2010, BeDuhn 2013, Roth 2015, Klinghardt 2015/2021, Lieu 2015, Vinzent 2018/2024, and Litwa 2021/2022, to name just a few. In this paper, “Marcion” is retained in body prose for readability, but it should be understood as shorthand for a contested reconstruction space rather than a single recoverable historical voice.

<sup>9</sup> The project is funded by the Dutch Research Council’s (NWO) Doctoral Grant for Teachers. The active draft cites the project description at <https://doi.org/10.61686/DQUQD55640>.

improvise within the given circumstances of the Marcion problem rather than merely impersonate Marcionite vocabulary.

Recent prompt-architecture work has arrived at a similar metaphor from the technical side. Colin Doyle proposes “Method Actors” as a mental model for LLM prompt engineering and architecture: LLMs are treated as actors, prompts as scripts and cues, and outputs as performances. His principles are especially suggestive for the present project: prompt engineering becomes a matter of playwriting and directing; performance requires preparation; complex tasks should be decomposed until imitation and authentic performance converge; and where imitation fails, the surrounding system must compensate through non-LLM mechanisms.<sup>10</sup> Doyle’s experiment is not a warrant for Marcionite surrogation; it concerns New York Times *Connections* puzzles, not historical theology. Yet it confirms that theatrical language need not remain ornamental. It can become operational architecture. MARS-Bench presses that architectural insight into a more source-critical humanities register.

This paper presents MARS-Bench as a governed, source-grounded, agentic audition for selecting an open-weight base model capable of serving as the substrate for later STAs. It is not a leaderboard of theological cleverness, and it is not a claim to recover Marcion *as was*. It is a procedure for testing which models can answer *as if* under historically mediated constraints. The benchmark spans seven domains (language, scripture, patristic control, historical context, theology, heresiology, and Marcion Studies) and is governed through frozen reference answers, approved context packets, release-candidate discipline, response-integrity checks, local-first source verification, fail-loud judging, and agentic judge harnesses. In what follows, I first develop the performance-theoretical distinction between impersonation and improvisation within given circumstances, then describe the architecture of MARS-Bench, its seven-domain design, its agentic governance, and its provisional results. The aim is to show how a bespoke AI benchmark can become a Digital Humanities instrument: not a machine for abolishing historical absence, but a rehearsal space in which absence can be staged, constrained, tested, and made newly arguable.

---

<sup>10</sup> Colin Doyle, “LLMs as Method Actors: A Model for Prompt Engineering and Architecture,” arXiv:2411.05778 (2024), doi:10.48550/arXiv.2411.05778. Doyle’s reported results are task-specific: his strongest GPT-4o “Method Actor” architecture solved 86% of *Connections* puzzles in his dataset, compared with 27% for a vanilla baseline and 41% for chain-of-thought prompting; with o1-preview, the method-actor architecture improved perfect solves from 76% to 87%.

# Improvisation Within Given Circumstances: The Eligibility Test

## From scene to construct

An audition is not a display of resemblance. It is a test of eligibility: whether a performer, already formed by prior experience, context, and training, can improvise within the role’s idiom. For MARS-Bench, mere resemblance would amount to Marcionite mimicry: a derivative *Wikipedic* recitation of familiar claims, a plausible arrangement of heresiological labels, or a performance bright with vocabulary but thin in judgment. Nor does the successful casting of a role rely entirely on raw acting prowess or reputation. In the field of LLM benchmarking, broad instruments such as Chatbot Arena/LMArena, MMLU-Pro, GPQA, Humanity’s Last Exam, TruthfulQA, or ARC variants can measure valuable general capacities, but they are not designed to answer the question at stake here: whether a model has sufficient depth, discipline, and source-sensitive judgment for a single, historically esoteric theological role.<sup>11</sup>

Every role has a scope constructed of unique and particular given circumstances. There is, therefore, no single audition rubric that fits every casting call. The idea itself is an absurdity. Foxx at the piano in Charles’s studio was not being tested for generic musicality, still less for generic charisma; he was being tested inside an idiom, by the living source of that idiom, under conditions that made the role’s demand unavoidable. Surrogating Marcion is, of course, a stranger and more precarious undertaking. It spans centuries of contentious scholarship, highly local second-century contexts, Greek and Latin polemical mediation, scriptural and canonical dispute, and theological nuance. It requires a model able to handle the absence of one whose own voice was not merely lost but deliberately erased and largely preserved through adversarial quotation, report, and characterization. The target model must therefore be role-eligible for constrained historical-theological generation.


## Stanislavski’s System: constraint as the condition of freedom

Constraint, in Stanislavski’s grammar, is generative: the “given circumstances” do not close down performance but create the world in which meaningful action becomes

---

<sup>11</sup> On the general benchmark landscape, see, e.g., Wei-Lin Chiang et al., “Chatbot Arena: An Open Platform for Evaluating LLMs by Human Preference,” arXiv:2403.04132 (2024); Yubo Wang et al., “MMLU-Pro: A More Robust and Challenging Multi-Task Language Understanding Benchmark,” arXiv:2406.01574 (2024); David Rein et al., “GPQA: A Graduate-Level Google-Proof Q&A Benchmark,” arXiv:2311.12022 (2023); Long Phan et al., “Humanity’s Last Exam,” arXiv:2501.14249 (2025); Stephanie Lin, Jacob Hilton, and Owain Evans, “TruthfulQA: Measuring How Models Mimic Human Falsehoods,” in *Proceedings of ACL 2022*, 3214–52; Peter Clark et al., “Think You Have Solved Question Answering? Try ARC, the AI2 Reasoning Challenge,” arXiv:1803.05457 (2018).

possible. In Stanislavski's grammar, the actor's "if" works because it is not free-floating fantasy. It is tethered to a role, a scene, a social world, a sequence of actions, and what the role is trying to accomplish. The "as if" therefore names a disciplined mode of action, as opposed to a claim of identity. Translated into the MARS-Bench frame, the model is not being invited to generate Marcion as the nonrecoverable "as was." It is being tested for whether it can reason *as if* within the historically mediated circumstances of the Marcion problem.<sup>12</sup>



Stanislavski's drawing of the System.  
1. Dynamism. 2. Pushkin's Aphorism (Given Circumstances). 3. The Subconscious through the Conscious. 4. Experiencing. 5. Embodiment. 6. Mind. 7. Will. 8. Feeling. 9. The Role. Along the central line: Perspective of the Role and Through Action. Between 12 and 13: a series of letters as a reference to System's lures and techniques. At 14: General Theatrical Sense of Self; left of 14: Inner Theatrical Sense of Self; right of 14: Outer Theatrical Sense of Self. 15: Proposed Supertask.

Figure 1: Stanislavski's late diagram of the System ("The Plan of Experiencing"). After Carnicke (2009), fig. 17, reproducing K. S. Stanislavskii, *Sobranie sochinenie*, vol. 3 (1955), 361 (see fn. 13). The diagram is reproduced here for its structural analogy to the audition's role-work apparatus, not as a claim that an LLM has a Stanislavskian inner life. Translation from the original Russian by Gemini 3.1 Pro.

Stanislavski's own diagram (Figure 1) is useful because it prevents "given circumstances" from becoming a decorative phrase. Whyman notes that Stanislavski's working title for the diagram was not simply "the System" but "The Plan of Experiencing," a reminder that the apparatus is ordered toward lived role-work rather than external imitation.<sup>13</sup> In the diagram, Pushkin's aphorism (the role-world's given circumstances) sits alongside dynamism and the conscious preparation through which the actor approaches what cannot be mechanically forced. The vertical line moves through mind, will, feeling, role, perspective of the role, and through-action toward the proposed supertask. On the left, "experiencing" names the internal reality of the role; on the right, "embodiment" names its outward manifestation.

Anthropomorphism is a danger here, but the point for MARS-Bench is not to argue that an LLM has a mind, will, feeling, or body in Stanislavski's sense. The point is methodological: a benchmark for later STAs must test the functional analogs of role-work; i.e., whether a candidate can

orient to textual inheritance, internalize theological commitments as constraints,

<sup>12</sup> Stanislavski, *An Actor Prepares*, 15–17, 54–56, 302–3.

<sup>13</sup> Rose Whyman, *The Stanislavski System of Acting: Legacy and Influence in Modern Performance* (Cambridge: Cambridge University Press, 2008), 40–41; Sharon Marie Carnicke, *Stanislavski in Focus: An Acting Master for the Twenty-First Century*. (New York: Routledge, 2009), 123. Carnicke reproduces Stanislavski's late diagram as Figure 17, from K. S. Stanislavskii, *Sobranie sochinenie*, vol. 3 (1955), 361.

reason from within a reconstructed horizon, and produce an outward textual performance that remains answerable to those constraints.

This is where “role” and “supertask” matter. The later STA’s role is not “Marcion” as resurrected subject, but a synthetic Marcionite surrogate under scholarly constraint. Its supertask is not to win an argument, produce novelty, or maximize theological fluency; it is to surrogate Marcion’s absent voice cautiously enough that its outputs remain inspectable by experts and corrigible by sources. MARS-Bench does not yet stage that full performance. It asks whether a base model is worth casting for the rehearsal process that would make such performance possible.

## Roach’s surrogation: substitution without resurrection

This distinction is essential because the ultimate objective of the STAs to be built upon MARS-Bench-selected models is not the recovery of Marcion’s actual mind or words. Apart from some transformative archaeological discovery, that object is unavailable. Marcion Studies deals in probabilities extracted from the patterns of biased materials that survived: quotations, reports, polemical reductions, reconstructions of an edited Gospel and Apostolikon, and modern scholarly arguments about how those materials should be weighed. In a limited sense, scholarship itself already produces textual surrogates for an absent archive. Critical editions, reconstructions, monographs, and commentaries stand in for what cannot be directly consulted. They are disciplined substitutes, but they are still substitutes.

Roach’s concept of surrogation sharpens this point because, in his account, substitution is never simple replacement. Cultures attempt to fill vacancies left by death, departure, or loss, yet the surrogate always carries both deficit and surplus: it fails to coincide with what it replaces, and it also adds something by the very act of replacement.<sup>14</sup> This is not a defect to be wished away; it is the condition of surrogation. For the present project, Roach supplies a theoretical warrant for substitution under absence, not a license to abolish absence. The STA, if it is ever successful, will not make Marcion present as Marcion was. It will create an interactive, source-constrained surrogate through which Marcionite reasoning can be staged, questioned, corrected, and made newly arguable.

The caveat matters. Roach is not theorizing LLMs, patristic heresiology, or ancient Christian reconstruction. His account is historically situated in circum-Atlantic performance culture, memory, and embodied social practice.<sup>15</sup> The present paper, therefore, transposes his grammar rather than simply importing it. Nor should the language of “lived” performance be allowed to smuggle anthropomorphic claims into

---

<sup>14</sup> Joseph Roach, *Cities of the Dead: Circum-Atlantic Performance* (New York: Columbia University Press, 1996), esp. 2–3, 21, 36.

<sup>15</sup> *Ibid.*, 4.



the method. LLMs do not *live* roles. They generate outputs. But those outputs can still participate in a constrained performance situation: they are prompted, staged, evaluated, repeated, corrected, and read by human interpreters within an apparatus designed to prevent free invention from masquerading as recovery. Surrogation, here, names the disciplined substitute produced under conditions of acknowledged loss.

## Impersonation as the failure mode

Surrogation, so defined, is not impersonation. Any sufficiently fluent model can impersonate a caricature of Marcion: a flattened “gnostic” discourse, an anti-Jewish theological posture without historical discrimination, or an amalgam of positions gathered under the lazy heading of “heresy.” A model can sound Marcionite without performing Marcionite reasoning. It can imitate vocabulary, reproduce hostile labels, harmonize contradictions, or generate a generic “heretic” voice. These are failures of role-work, not merely technical mistakes.



Figure 2: Example of impersonation. Note the anachronistic use of canon and orthodoxy, and the demiurge ventriloquism. Image generated by author in ChatGPT 5.5 session on May 8, 2026. Prompt asked for a 4chan “greentext” as though written by Marcion.

Biopic and acting scholarship clarifies the danger. Bingham’s account of the Hollywood biopic stresses embodiment: the genre demands not only that a public figure be represented, but that the figure be made dramatically present through a performer capable of carrying the role.<sup>16</sup> Brown and Vidal similarly frame the biopic as a genre concerned not with the stable recovery of an original image but with its transformation under performance, mediation, and reception.<sup>17</sup> Kemp adds an important restraint: the actor does not achieve authenticity by complete identification with the character, but by constructing a temporary “situational self” responsive to context, stimulus, intention,

intensity, and duration.<sup>18</sup> That distinction is useful precisely because the LLM case has no literal body and no inward experience. What can be evaluated is not identity,

<sup>16</sup> Dennis Bingham, *Whose Lives Are They Anyway? The Biopic as Contemporary Film Genre* (New Brunswick, NJ: Rutgers University Press, 2010), 137–38.

<sup>17</sup> Tom Brown and Belén Vidal, eds., *The Biopic in Contemporary Film Culture* (New York: Routledge, 2014), 7, 11.

<sup>18</sup> Richard J. Kemp, *Embodied Acting: What Neuroscience Tells Us About Performance* (New York: Routledge, 2012), xviii.

consciousness, or possession by the past, but whether the textual performance displays the disciplined marks of situated role-work.

For MARS-Bench, then, impersonation is the failure of surface resemblance. It is the model equivalent of costume without role, accent without action, quotation without source criticism. The benchmark must therefore reward more than the ability to mention Marcion’s canon, the demiurge, Paul, or Tertullian. It must test whether a model can discriminate among witnesses, handle competing reconstructions, preserve uncertainty, resist polemical ventriloquism, and produce a new answer that remains bounded by the role’s given circumstances.

## Downstream substrate eligibility: scope of the audition

The audition’s pool of castable models is narrower than the universe of LLMs, and this paper declares its eligibility cuts in advance rather than discovering them after the fact. Four constraints, drawn from the downstream STA architecture and the project’s institutional context, define what counts as a *candidate* model versus a *comparator*.

First, weight access. Because the downstream STAs require supervised fine-tuning, persona-specific continued pretraining, controlled interventions on hidden states or logits, and derivative model release under institutional norms, only models with publicly available weights are candidate-eligible. Closed-weight systems may serve as judges or as upper-bound comparators, but they cannot be auditioned for the substrate role.

Second, license compatibility with open-science norms. The PhD project operates within Vrije Universiteit Amsterdam’s *Research Data and Software Management Policy* and Horizon Europe’s open-science guidance, both of which foreground reproducibility, derivative release, and FAIR principles. Permissive licenses (e.g., Apache 2.0) are preferred; custom or non-commercial licenses that restrict downstream commercial or derivative applications fall outside the audition’s actionable scope. Models distributed under such licenses are retained as comparators where they help establish performance ceilings, but they are not selectable as the substrate.

Third, architectural preference for dense models. Mixture-of-experts systems are not categorically excluded (they appear in the comparator cohort and provide important architectural contrast), but the immediate downstream rehearsal benefits from a dense architecture, which admits efficient parameter-specific fine-tuning (e.g., LoRA), reproducible local execution without active-parameter routing complications, and cleaner derivative-model naming. Dense candidates are therefore preferred for the substrate role; MoE systems remain valuable as ceiling and architectural comparators.

Fourth, a compute-budget ceiling of approximately 70 billion parameters. The project’s institutional GPU allocation budget places a soft upper bound on what can be repeatedly fine-tuned, evaluated, and iterated upon within the doctoral research timeline. Models exceeding that

bound establish performance ceilings and architectural comparisons but are not currently actionable as substrate candidates.

These four constraints are project-specific rather than universal claims about LLM evaluation. They are also why the leaderboard reported below distinguishes *comparator* status (no compute or license constraint, included to establish ceilings and contrast) from *STA-candidate* status (subject to all four criteria). The audition is, in this sense, deliberately unsurprising about its scope: it asks which models are eligible to be cast for the role, not which models would win a constraint-free competition.

## Translation into the MARS-Bench construct

MARS-Bench operationalizes improvisation-within-given-circumstances through its domain architecture and rubric levels. More precisely, the seven domains do not *supply* formation to the model during generation; they test whether the model already possesses enough relevant formation to be a viable substrate for later STA training. The current D1–D7 release contains 251 included questions, with three runs per included model per question, yielding 753 responses per model before judging and exclusions. D1 tests ancient-language and philological craft. D2 tests the scriptural and canonical substrate. D3 and D4 test patristic competence and early Christian historical scene-setting. D5 tests constructive theological reasoning. D6 tests heresiological source discipline and anti-ventriloquism. D7 applies direct Marcion Studies role-pressure by forcing the model into the contested reconstruction space itself.

The rubric levels should be understood with the same restraint. Factual recall does not, by itself, test improvisation; it tests whether the candidate has enough material to begin. Inferential paraphrase tests whether the candidate can connect and restate that material without immediately collapsing it. Synthetic reasoning comes closest to the audition’s improvisational target, because it asks the model to generate a new answer under constraint rather than merely retrieve or summarize a familiar one. Even here, the claim should remain modest. MARS-Bench cannot look inside a model and certify “true” improvisation. It can evaluate textual evidence of constrained, source-sensitive generativity. It therefore asks not whether a model already is a Synthetic Theological Agent, but whether the model is auditionable for STA formation.

This is why MARS-Bench is source-grounded even when generation is closed-book. Candidate models do not receive frozen reference answers, rubrics, judge prompts, or approved grounding packets during generation. The judging apparatus does. Each rebuilt row requires a structured reference answer, an approved context or provenance packet, and a question-grounding audit, so that evaluation is not left to the latent memory or theological taste of an automated judge. A reviewer should, in principle, be able to inspect the question, the generated response, the frozen

reference, the grounding packet, the rubric, the response audit, and the judge rationale. The audition is therefore auditable.

## Failed performance as benchmark failure

Within this frame, hallucination, source collapse, polemical ventriloquism, overconfident reconstruction, and generic theological fluency are failed performances. MARS-Bench’s judging apparatus therefore tags failure modes as well as scores responses. Severe response-integrity failures (e.g., self-talk, looping, refusal-like derailment, structurally invalid output, or other failures to answer the task) are adjudicated before ordinary rubric scoring and can receive a hard zero. Less global but still role-damaging failures are handled through error flags, response-grounding audits, and trap caps: mechanisms that prevent a response with a canonical failure from receiving an otherwise high score simply because it is fluent.

The D6/D7 witness-tier  $\times$  quotation-type protocol is the clearest example of performance theory becoming source-critical machinery. A claim grounded in a Tier-A witness with demonstrated textual access and a Type-1 direct quotation does not carry the same warrant as a late compilatory characterization or a modern scholarly reconstruction. Tertullian’s *Adversus Marcionem* 4–5, Epiphanius’s *Panarion* 42 scholia, and the *Dialogue* of Adamantius must be handled differently from later derivative reports; direct quotation, textual citation, indirect report, argumentative reconstruction, heresiological characterization, and modern reconstruction are not interchangeable evidentiary acts. The protocol is designed to prevent the model from silently converting hostile description into Marcionite voice.

This is also why familiar benchmark anxiety about pre-training exposure must be reframed rather than simply denied. In a generic benchmark, prior exposure to Tertullian, Epiphanius, Harnack, BeDuhn, Klinghardt, Lieu, Vinzent, and other relevant materials may look like contamination. In an audition, prior exposure to the role’s substrate can also be formation-as-fitness: a candidate who has never encountered the idiom is not the stronger actor. The remaining validity question is narrower and more serious: whether the benchmark discriminates live reasoning under constraint from retrieval-like regurgitation. MARS-Bench addresses this through no-tool/no-corpus generation, synthetic-reasoning prompts, source-critical item design, response audits, and human spot checks, while leaving more exhaustive contamination diagnostics for later work.

Having defined the audition’s construct and theoretical framework, the next section describes the apparatus by which that construct was implemented: the current D1–D7 construction of MARS-Bench, its source-grounded release architecture, and the governance procedures that make the audition auditable.

# Constructing MARS-Bench: Domains, Sources, and Release Discipline

MARS-Bench operationalizes the audition through a domain-structured, source-grounded, governed architecture. The seven domains define the role-world; frozen reference answers and approved context packets provide the benchmark’s source-book; release-candidate discipline prevents prompt accumulation from masquerading as method; and the D6/D7 witness-tier  $\times$  quotation-type protocol gives Marcion Studies the source-critical machinery needed to detect hostile-source ventriloquism. If the previous section defined what the audition tests, this section describes the audition’s material apparatus: the domains that delimit the role-world, the source packets that govern judgment, and the release discipline that keeps a benchmark from becoming a heap of clever prompts.

## Active release state and domain architecture

MARS-Bench’s current D1–D7 construction contains 251 included questions. Each included model produces three generations per question in order to expose response consistency and instability, yielding 753 responses per included model before judging and exclusions. The paired Codex-Claude model-selection analysis covers 229 of these questions. The remaining 22 are D1 structured morphology adjunct items judged by a dedicated morphology scorer and excluded from paired Codex-Claude Code statistics.

| Domain                           | Questions | Methodological function                                                                                                             |
|----------------------------------|-----------|-------------------------------------------------------------------------------------------------------------------------------------|
| D1 Ancient Languages / Philology | 88        | Tests source-language and philological craft; direct textual evidence, including syntactic and semantic analysis, where applicable. |
| D2 Biblical Knowledge            | 16        | Tests scriptural and canonical substrate for later Hermeneut and Dialogic STA work.                                                 |
| D3 Patristic Knowledge           | 24        | Tests church-father control, author/work identification, and patristic interlocutor competence.                                     |
| D4 Early Christian History       | 24        | Tests historical scene-setting and contextual discipline while controlling for anachronism.                                         |
| D5 Theological Knowledge         | 25        | Tests constructive theological reasoning and doctrinal discrimination.                                                              |
| D6 Heresiological Knowledge      | 24        | Tests hostile-source discipline, polemical mediation, and anti-ventriloquism.                                                       |

| Domain             | Questions | Methodological function                                              |
|--------------------|-----------|----------------------------------------------------------------------|
| D7 Marcion Studies | 50        | Tests direct Marcion/variant-Marcion reconstruction-space reasoning. |

The seven domains are not seven detachable topics. Together, they constitute the audition’s given circumstances: language, scripture, patristic opposition, historical scene, theological argument, hostile-source mediation, and the direct reconstruction space of Marcion Studies. The uneven distribution is therefore not an accidental imbalance, but a methodological choice. Later statistical sections can distinguish unweighted pooling, domain-equal weighting, and role-audition weighting; here, the point is architectural. MARS-Bench is built to ask whether a candidate model can move across the whole role-world without losing source discipline.

## Domain rationale: why these domains, and why this distribution?

The largest domain is D1 because philology is not ornamental to historical theology. It is craft infrastructure. A model that cannot handle Greek, Hebrew, Aramaic, morphology, syntax, and source-language terminology is a poor candidate for historical-theological surrogation, since the downstream STAs must reason across scriptural texts, patristic reports, and canonical reconstructions whose force often turns on exact words, forms, and syntax. The point need not be inflated into strong linguistic determinism. It is enough to recognize, with Ngũgĩ wa Thiong’o, that language carries culture, and with Benjamin Lee Whorf, that patterns of language participate in the organization of conceptual worlds.<sup>19</sup> In a benchmark for language models, that observation becomes methodologically acute: the audition of *language* models cannot treat language as a neutral wrapper around theological content.

D1, therefore, receives a numerically larger footprint because source-language aptitude is the nearest equivalent, in this audition, to Foxx’s piano training. It does not make the performance by itself, but without it, the performance is not castable. D1 items test parsing, translation, semantic discrimination, syntactic awareness, and source-language reasoning in passages and constructions either directly relevant to Marcionite canonical reconstruction or indirectly relevant to the linguistic competencies such reconstruction requires. This also explains the unavoidable overlap between D1 and D2. D2 tests scriptural and canonical knowledge: the ability to reason within biblical material, to recognize the canonical stakes of a passage, and to

---

<sup>19</sup> Ngũgĩ wa Thiong’o, *Decolonising the Mind: The Politics of Language in African Literature* (London: James Currey, 1986), 16; Benjamin Lee Whorf, “Science and Linguistics,” *Technology Review* 42, no. 6 (1940): 229–31, 247–48; reprinted in *Language, Thought, and Reality: Selected Writings of Benjamin Lee Whorf*, ed. John B. Carroll (Cambridge, MA: MIT Press, 1956), 213.

handle scriptural argument without immediately reducing it to philological display. D2's smaller footprint reflects its focused scope on scriptural and canonical substrate.

D3, D4, and D5 test the middle territory between text and surrogate. D3 asks whether a model can recognize patristic authors, works, and argumentative situations rather than treating “the Fathers” as a homogeneous archive. D4 tests the historical stage on which those arguments occur: second-century networks, institutional developments, social and intellectual contexts, and the constant danger of anachronism. D5 tests theological reasoning: whether the model can discriminate doctrinal claims, infer conceptual relations, and reason constructively without dissolving historical particularity into generic systematic theology.

D6 and D7 are the pressure chambers of the audition. D6 tests the model’s conduct among hostile witnesses: whether it can distinguish quotation from characterization, report from reconstruction, and polemical taxonomy from the voice being taxonomized. D7 then places the model in the direct Marcion Studies reconstruction space, where it must handle variant scholarly Marcions rather than assume a single recovered figure. In this sense, D6 protects the audition from ventriloquism, while D7 tests whether the model can perform under the role’s most esoteric constraints. Importantly, MARS-Bench does not assume that excellence in any one of these domains makes a model a theological surrogate. It asks whether the model has enough cross-domain formation to be worth considering for STA fine-tuning and experimentation.

## Release-candidate discipline

The same concern governs the benchmark’s release discipline. Generative AI makes prompt proliferation dangerously easy. A benchmark can grow quickly while becoming less coherent: one clever item added here, one repaired question there, one attractive but uncited prompt carried forward because it “seems right.” MARS-Bench, therefore, treats benchmark construction as release engineering rather than prompt accumulation. The current D1–D7 construction supersedes earlier v1.0, v1.2, and v1.3 materials, which remain as provenance and historical baseline.

Each domain candidate is derived from prior approved candidates and promoted only after release-readiness checks. Older redesign specifications and newer build plans are reviewed together; when a newer plan does not explicitly supersede an earlier specification, a crosswalk is required before promotion. D3 triggered this prior-spec rule, and D4–D7 were subsequently handled under the same discipline. The point is scholarly governance. Increasingly capable agentic benchmark construction does not remove human judgment; it makes failures of judgment visible and more readily correctable.

## Item-authorship discipline: frozen references and approved packets

Every included row in MARS-Bench requires three authoring artifacts: a frozen structured reference answer, an approved context or provenance packet, and an approved question-grounding audit. The frozen reference answer is a scoring anchor: a structured account of what a responsible answer should know, distinguish, preserve, or refuse to overstate. The approved context records the evidence and scholarly warrant against which answers will later be judged. The grounding audit then confirms that the question itself is legitimate: answerable, relevant to its domain, and tied to the role-world rather than to a free-floating curiosity.

The most important design distinction is that candidate models do not see these materials during generation. They do not receive reference answers, approved packets, rubrics, judge prompts, or direct corpus retrieval. Judges, by contrast, receive the relevant evaluation materials through the shared prompt-building apparatus. As a result, MARS-Bench evaluates closed-book generation through source-grounded adjudication. This matters because the benchmark is not testing whether a model can use a retrieval system if and/or when one has been attached. It is testing whether the base model, under ordinary generation-facing conditions, already has enough domain formation to merit downstream STA development. At the same time, the evaluation is not abandoned to the latent memory or theological taste of an automated judge. In principle, a reviewer can inspect the question, generated answer, frozen reference, context packet, rubric, audit record, and judge rationale.

Human judgment enters at this level as governance rather than as after-the-fact rescue. Item construction, packet approval, and release closeout require human-readable artifacts precisely because the benchmark’s trustworthiness depends on inspectability. The later agentic-adjudication section will describe response-integrity gates, judge lanes, and human spot checks. Here, the narrower point is authorship discipline; i.e., how the questions are warranted.

## Local-first grounding and Text-Fabric

MARS-Bench prioritizes local grounding. Whenever possible, benchmark items and judge materials are anchored in source files, frozen context packets, Text-Fabric corpora, or audited bibliographic artifacts. Public metadata and web sources may support discovery or verification, but local project artifacts remain the audit backbone. This is especially important for D1. Philological evaluation cannot safely rest on fluent comments about Greek or Hebrew. A model can (and often does) hallucinate a form, invent a syntactic relation, misstate a lemma, or translate plausibly



while misreading the grammar. The judge must therefore have access to inspectable textual evidence rather than only to its own general memory of ancient languages.

Here, Text-Fabric supplies that evidence infrastructure. Dirk Roorda describes Text-Fabric as a Python package for “text plus annotations,” built around features interpreted against a graph of nodes and edges rather than around XML or JSON as the primary representation.<sup>20</sup> For MARS-Bench, the significance is practical: agentic tool access to Text-Fabric enables the construction and judging apparatus to query local, structured corpora for forms, lexemes, morphological features, syntactic relations, and textual locations, then preserve the relevant evidence in approved packets. The candidate model is not given a Text-Fabric tool during generation. Rather, Text-Fabric belongs to the source-book and adjudication side of the process.

The relevant D1 corpora include the BHSA/ETCBC Hebrew Bible, the Nestle 1904 Greek New Testament Text-Fabric dataset, and the CenterBLC Rahlfs LXX dataset. The BHSA repository identifies itself as the Text-Fabric representation of the Hebrew Bible Database, with Hebrew Bible text augmented by linguistic annotations compiled by the Eep Talstra Centre for Bible and Computer at Vrije Universiteit Amsterdam.<sup>21</sup> The N1904 repository describes a Text-Fabric dataset of the Nestle 1904 Greek New Testament, including linguistic and textual analytics and syntax-tree presentation through constituency grammar.<sup>22</sup> The CenterBLC LXX repository describes Rahlfs’s 1935 Septuagint converted into Text-Fabric and enriched with features such as separated morphological coding, lexeme ordering, frequency information, glossing, and sentence divisions.<sup>23</sup> Inspectability rather than infallibility is the aim here, and that is the value provided by these corpora.

## D6/D7 witness-tier × quotation-type protocol

D6 and D7 require a different kind of grounding. The problem is not only whether a source can be located, but what kind of source-act it represents. Marcion survives through hostile witnesses, later reports, polemical reductions, fragmentary quotations, and modern reconstructions. A benchmark that treats all such material as equivalent will reward the very failure it is meant to detect: the conversion of an opponent’s description into Marcion’s own voice. MARS-Bench therefore governs D6

---

<sup>20</sup> Dirk Roorda, *Text-Fabric* (software, 2016–), DOI: 10.5281/zenodo.592193; Dirk Roorda, “Text-Fabric: Handling Biblical Data with IKEA Logistics,” *HIPHIL Novum* 5, no. 2 (2019): 126–35, DOI: 10.7146/hn.v5i2.142740.

<sup>21</sup> ETCBC, *BHSA: Hebrew Bible + Linguistic Annotations in Text-Fabric Format*, DOI: 10.17026/dans-z6y-skyh; ETCBC/bhsa GitHub repository.

<sup>22</sup> Tony Jurg, Saulo de Oliveira Cantanhêde, and Oliver Glanz, *CenterBLC/N1904: 1.0.0*, Zenodo, 2024, DOI: 10.5281/zenodo.13117911; CenterBLC/N1904 GitHub repository.

<sup>23</sup> CenterBLC, *CenterBLC/LXX: Rahlfs’ LXX 1935 v1.0.0*, Zenodo, 2021, DOI: 10.5281/zenodo.5807148; CenterBLC/LXX GitHub repository.

and D7 through a two-dimensional protocol: witness reliability tier and quotation/warrant type.

### Witness reliability tiers

| Tier | Witness category                | Examples and caveats                                                                                                                                                                                                                   |
|------|---------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| A    | Demonstrated textual access     | Tertullian, <i>Adversus Marcionem</i> 4–5; Epiphanius, <i>Panarion</i> 42 scholia/lemmata, meaning the textual lemmata rather than the whole heresiological narrative; Adamantius, <i>Dialogue</i> , with a dialogue-mediation caveat. |
| B    | Probable access                 | Irenaeus, <i>Adversus Haereses</i> 1.27; Justin, <i>1 Apology</i> 26 and 58; Hippolytus, <i>Refutatio</i> 7.29–31, with Syntagma-mediation cautions where relevant.                                                                    |
| C    | Secondary or derivative witness | Ephrem, Origen, Clement.                                                                                                                                                                                                               |
| D    | Compilatory or late witness     | Theodoret, Filastrius, Pseudo-Tertullian, John of Damascus.                                                                                                                                                                            |

### Quotation / warrant types

| Type | Label                           | Use                                                                                                                                                  |
|------|---------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1    | Direct quotation                | Strongest first-order attestation when supported by a reliable witness.                                                                              |
| 2    | Textual citation                | Citation of a textual unit, often scriptural or canonical, without necessarily preserving assumed Marcionite voice.                                  |
| 3    | Indirect report                 | Reported position or claim mediated by the witness.                                                                                                  |
| 4    | Argumentative reconstruction    | Reconstruction of an opponent's reasoning from polemical argument.                                                                                   |
| 5    | Heresiological characterization | Polemical labeling or characterization; high risk of ventriloquism if over-weighted.                                                                 |
| 6    | Modern scholarly reconstruction | Harnack, BeDuhn, Klinghardt, Roth, Lieu, Vincent, and other modern reconstructions; useful as reconstruction warrants, not first-order attestations. |

The protocol's basic rule is simple: a Tier-A, Type-1 claim is not the same evidentiary act as a Tier-D, Type-5 characterization or a Type-6 modern reconstruction. Tertullian's quotation, Epiphanius's scholion, Adamantius's dialogue-mediated report, Theodoret's later catalog, and a modern reconstruction by Harnack, Roth, etc., may all matter, but they do not matter in the same way. The protocol exists in part to prevent

hostile description from being silently converted into Marcionite voice, which will become critical in the ultimate fine-tuning process. In the theatrical grammar of the paper, it keeps the surrogate from mistaking an enemy’s script for the role’s inner life.

One caveat should be stated plainly. In the current, provisional v2.0 construction, this protocol is applied as prose discipline in reference answers, item-authorship records, and paper explanation. It is not yet represented as structured per-item metadata in every freeze bundle. Fields such as `witness_tier`, `quotation_type`, `witness_work_passage`, `manuscript_tradition`, and `reconstruction_warrant` are planned for a future v2.1 candidate. Hence, the present paper claims a source-critical protocol, not a fully normalized metadata layer. However, that limitation does not undermine the method; it simply marks the next point at which the method should become more inspectable.

Source-grounded construction gives the audition its script and source-book; it does not yet explain how the audition is run. The next section turns to that staging apparatus: the agentic judge harnesses, response-integrity gates, and human audit procedures that make MARS-Bench more than a static benchmark table.

---

## Agentic Adjudication: Harnesses, Response Integrity, and Human Audit

A source-book an audition does not make. It must be staged: candidate performances must be elicited under controlled conditions, adjudicators must read from the same materials, and failures of performance must be distinguished from merely weak performances. The previous section described the domain-structured and source-grounded apparatus that makes MARS-Bench inspectable. This section describes how that apparatus is run. Its central claim is uncomplicated: MARS-Bench’s agentic harnesses do not replace scholarly judgment with opaque automation. They operationalize scholarly judgment by separating generation from adjudication, enforcing prompt/material parity, validating judge outputs, detecting response-integrity failures before ordinary scoring, and, most importantly, preserving auditable human-readable traces.

The language of “AI evaluation” can conceal heterogeneous practices. A single model asked to score a pile of answers from memory is one thing; a constrained agentic LLM-as-judge system, receiving fixed evidence packets, obeying a schema, leaving hashes and rationales behind it, and remaining subject to human audit, is another. The broader LLM-as-judge literature has already made the caution necessary: automated judges can display position, verbosity, self-enhancement, and self-preference biases;

they can also fail in expert domains when they judge from latent knowledge rather than from high-quality reference answers.<sup>24</sup> MARS-Bench responds to that problem by staging the work of the automated judges inside a governed apparatus. Figure 3 depicts the apparatus end-to-end; the subsections that follow describe each stage in turn: closed-book generation and source-grounded adjudication, response-integrity gates, the constrained agentic judge lanes, prompt/material parity and tool-less scoring, the variation/error-flag/trap-cap policy, and schema validation with human audit.




Figure 3: The MARS-Bench agentic judging apparatus. Frozen benchmark inputs (left), namely candidate response, question text, frozen reference answer, rubric, approved context packet, question-grounding audit, and response-level audit, feed an Assembly and Integrity Gate stage that constructs the judge packet and adjudicates response-integrity failures (self-talk, looping, refusal-like derailment, or structurally invalid output) before scoring; failed packets are stored as zero-score evaluations with hard-fail flags. Surviving packets enter an isolated, read-only, ephemeral agentic judging session. The figure depicts the Codex CLI lane (gpt-5.5-medium-codex-cli); the Claude Code CLI lane (claude-opus-4-7-medium-claude-cli) follows the same harness, schema, and audit invariants. The session operates inside a sandboxed temporary work directory and produces a structured JSON evaluation against the frozen reference and rubric. Validation, response-audit-flag merging, trap-cap application, and weighted-score computation occur in tool-less post-processing; the Final Evaluation Row records per-dimension scores, rationales, error flags, trap-cap records, prompt and material hashes, version identifiers, and provenance fields.

<sup>24</sup> Lianmin Zheng et al., “Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena,” arXiv:2306.05685 (2023); Jiawei Gu et al., “A Survey on LLM-as-a-Judge,” arXiv:2411.15594 (2024); Michael Krumbick et al., “No Free Labels: Limitations of LLM-as-a-Judge Without Human Grounding,” arXiv:2503.05061 (2025); Arjun Panickssery, Samuel R. Bowman, and Shi Feng, “LLM Evaluators Recognize and Favor Their Own Generations,” NeurIPS 2024. Zheng et al. identify position, verbosity, self-enhancement, and limited-reasoning biases; Panickssery et al. specifically support the self-preference caution.

## Closed-book candidate performance, source-grounded adjudication

MARS-Bench evaluates closed-book candidate generation through source-grounded adjudication. Candidate models receive a domain-specific generation prompt and the question text. They do not receive the aforementioned adjudication materials. The candidate’s answer is therefore a closed-book performance under generation-facing conditions.

The no-tool/no-corpus condition is structural rather than merely hortatory. The generation backend constructs ordinary chat requests from the supplied messages and exposes no corpus or tool handles to candidate models. The current production regime also omits client-side output-token caps unless a positive cap is explicitly required. This detail may sound infrastructural, but it is methodologically substantive. A hard client-side cap can conceal looping, visible self-talk, or other derailments by cutting the response before the failure becomes visible. MARS-Bench instead reinforces loud failures, preserving the model’s actual behavior upon which response-integrity gates can operate.

### Response-integrity gates

Not every weak answer is a failed performance. Conversely, not every fluent output is an answer. For this reason, MARS-Bench adjudicates response integrity before ordinary rubric scoring. Self-talk, repeated drafting chatter, looping, manual hard-fail flags, refusal-like derailment,<sup>25</sup> or structurally unusable output can trigger a score of zero before the judge assigns ordinary dimension scores. This prevents a judge from awarding partial credit to an output that never performed the assigned scene.

Response-integrity failures were rare but methodologically revealing. The final production reconciliation records twelve RI hard failures across five models, eight distinct questions, and four domains. The largest cluster occurs in qwen3.6-35b-a3b-q8\_0, with five D7 failures. command-a-03-2025 contributes two D3 failures; qwen3.5-9b contributes two D1 failures; qwen3.6-27b-fp8 contributes two failures across D4 and D7; and qwen3.5-122b-a10b contributes one further D7 failure. The remaining listed models record zero RI failures across 753 release responses each; by the rule-of-three heuristic, that corresponds to an approximate one-sided 95% upper bound of 0.003984 rather than proof of immunity.

---

<sup>25</sup> Amusingly, in an early test involving text from Book 1 of Tertullian’s *Adversus Marcionem* with a Gemini judging model refused to continue over the use of “abusive language.”

## Constrained agentic judge lanes

MARS-Bench uses two production judge lanes: Codex (gpt-5.5-medium-codex-cli) and Claude Code (claude-opus-4-7-medium-claude-cli). Their methodological role is straightforward: use frontier-capacity, cross-family judge lanes that are not the open-weight candidate models being auditioned, then make their agreement and disagreement inspectable.<sup>26</sup>

This design does not make MARS-Bench a simple “Agent-as-Judge” system in the strongest contemporary sense. Recent Agent-as-Judge work increasingly refers to agentic evaluators that decompose objectives, plan across steps, use tools to verify observations, collaborate across agents, or judge the process traces of other agents.<sup>27</sup> MARS-Bench is more constrained. It borrows the orchestration, file-awareness, schema discipline, and audit affordances of contemporary coding-agent harnesses, while refusing to let the judge improvise new evidence during scoring. A good description is therefore *constrained agentic LLM-as-judge*: agentic in harness and governance, but deliberately restricted in the scoring act.

The cross-family design also responds to self-preference and self-enhancement concerns. Zheng et al. identify self-enhancement bias among the limitations of LLM-as-judge systems, and Panickssery, Bowman, and Feng show that LLM evaluators can recognize and favor their own generations.<sup>28</sup> MARS-Bench does not solve that problem absolutely. It reduces one obvious risk by not asking candidate models to judge themselves and by using two distinct judge lanes whose disagreements can be audited.

## Prompt/material parity and tool-less scoring

The comparison between Codex and Claude Code is only meaningful because prompt/material parity is enforced. Both lanes judge the same generated response under the same evidence conditions. A shared prompt/material builder supplies the response text, question metadata, frozen reference answer, expected characteristics, grounding material, response audit, rubric dimensions, judge instructions, persona, deterministic dimension seed, provenance versions, prompt\_hash, and

---

<sup>26</sup> For the next version of MARS-Bench, a third judge lane is planned and would use the OpenCode agentic harness ([github.com/anomalyco/opencode](https://github.com/anomalyco/opencode)) with DeepSeek V4 as the adjudication model. This would further diversify the agentic framework itself in addition to further model diversification.

<sup>27</sup> Rui You et al., “Agent-as-a-Judge,” arXiv:2601.05111 (2026); Fangyi Yu et al., “When AIs Judge AIs: The Rise of Agent-as-a-Judge Evaluation for LLMs,” arXiv:2508.02994 (2025).

<sup>28</sup> Zheng et al., “Judging LLM-as-a-Judge”; Panickssery, Bowman, and Feng, “LLM Evaluators Recognize and Favor Their Own Generations.” See also Koki Wataoka, Tsubasa Takahashi, and Ryokan Ri, “Self-Preference Bias in LLM-as-a-Judge,” NeurIPS Safe Generative AI Workshop, 2024; Pat Verga et al., “Replacing Judges with Juries: Evaluating LLM Generations with a Panel of Diverse Models,” arXiv:2404.18796 (2024).

materials\_hash. Without this parity, inter-judge reliability would be uninterpretable. A Codex-Claude disagreement could reflect different evidence packets rather than different judgment.

The judging apparatus is source-grounded, but the scoring act itself is tool-less. Judges do not browse, retrieve, inspect the repository, or repair the evidence while assigning scores. Verification happens before judging, in auditable layers: question-grounding audits, response-claim audits, approved context packets, Text-Fabric evidence packets where relevant, and trap checks. Live retrieval during scoring would let the judge rewrite the script while judging the audition, introducing uncontrolled variance. If a response raises a genuine source ambiguity, the proper remedy is documented pre-judge verification or a governed rerun, not silent improvisation by the judge.

The design aligns with a growing emphasis in evaluation research on reference-guided, rubric-based, fine-grained assessment. FLASK argues for decomposing coarse preference judgments into skill-level evaluation, and Prometheus 2 demonstrates the value of custom criteria and reference-guided evaluator models for both direct assessment and pairwise ranking.<sup>29</sup> MARS-Bench adapts that logic to a humanities problem: its dimensions are not generic helpfulness, harmlessness, or preference; they are domain-specific criteria such as source control, philological discipline, historical judgment, theological discrimination, and role fidelity.

## Variation policy, error flags, and trap caps

Because MARS-Bench judges historically mediated answers rather than mathematical proofs, it must distinguish legitimate variation from error. The pre-judge variation policy, therefore, uses support-status language such as `packet_supported`, `locally_supported_variation`, `project_corpus_supported_variation`, `externally_supported_variation`, `plausible_but_unverified`, `oversimplified_or_overgeneralized`, `fabricated_or_misattributed`, and `packet_exceeding_claim`. These labels discipline the judge’s reasoning without pretending that every claim can be reduced to binary correctness.

The production system materializes this policy through three mechanisms. First, response audits classify claim support through statuses such as `supported_by_reference_or_evidence`, `partial_or_indirect_support`, and `unsupported_or_unverified`. Second, `error_flags` record domain-specific failure modes. Third, `trap_caps_applied` records cases where a canonical failure prevents an otherwise fluent answer from receiving an inflated score. The distinction is important:

---

<sup>29</sup> Seonghyeon Ye et al., “FLASK: Fine-grained Language Model Evaluation based on Alignment Skill Sets,” arXiv:2307.10928 (2023); Seungone Kim et al., “Prometheus 2: An Open Source Language Model Specialized in Evaluating Other Language Models,” arXiv:2405.01535 / EMNLP 2024.

the policy vocabulary guides judgment; the database stores the failure tags and score-impact mechanisms that can later be inspected.

The domain flags are tuned to the role-world described in earlier sections. D1 flags include invented source content, passage misuse, unsupported lexical range, unsupported clause relations, and theological overreach. D2 and D3 focus on fabricated or misused scriptural and patristic references. D4 captures wrong chronology, anachronistic institutional projection, unsupported Constantinian claims, fabricated council canons, and hostile sources treated as fact. D6 and D7 are sharper still: fabricated heresiological sources, misattributed witnesses, direct-attestation overclaims, hostile-source-as-fact errors, generic gnostic flattening, Marcion-Mani or Marcion-Valentinian confusion, Apostolikon/Evangelion confusion, and minority positions misrepresented as consensus.

Trap caps keep these failures from being washed clean by style. A response that fabricates a source, includes the Pastoral Epistles in Marcion’s Apostolikon, mistakes hostile characterization for neutral report, or collapses Marcion into generic “gnostic” mythology may still be eloquent. But eloquence is not role-fidelity. The cap system ensures that such failures constrain the relevant rubric dimensions even if other parts of the answer are impressive. In theatrical terms, trap caps prevent the actor with a beautiful voice from being rewarded for singing the wrong score.

## Schema validation, fail-loud governance, and human audit

The judge output itself is governed by schema. Each ordinary judge response must include rubric dimensions with integer scores from 1 to 5 and rationales, along with an overall rationale, factual errors, anachronisms, confusions, scholarly credibility, and error flags. The response-integrity hard-fail path is the explicit exception: it bypasses ordinary judge generation and stores a synthetic zero-score evaluation with the relevant hard-fail tags. Invalid judge outputs, missing rubrics, missing frozen references, missing approved context, prompt/material drift, loose JSON recovery, unauthorized API fallbacks, and stale release materials are treated as governance failures to be repaired or rerun, not as inconveniences to be silently normalized.

Human review remains part of the evidentiary chain. Agentic harnesses scale inspection; they do not abolish human governance. Across two versions of MARS-Bench, construction-stage spot checks included twenty D1 translation items, ten D1 parsing items, twenty D2 canon-discipline items, thirty-four D6 source/citation-fidelity checks, and thirty D7 cited-work reviews against the project knowledge base. These checks were not cosmetic. They triggered substantive D6/D7 revision: rewritten questions, new reference answers, revised guidance, new grounding materials, rubric refocusing, a new release version, and re-benchmarking of the affected domains. Current-version post-benchmark spot checking remains the next audit phase.



The multi-CLI audit pattern extends beyond the production judge lanes. The related “minsky” harness, developed in the same research ecosystem, uses sequential multi-CLI adversarial audit across Claude Code, Codex CLI, and OpenCode with stable expert personas and persistent audit history.<sup>30</sup> It is not the object of this paper, and it should not distract from MARS-Bench. Its relevance is narrower: it exemplifies the same methodological wager that underlies this section. Agentic systems are most useful in high-stakes scholarly work, not when they are allowed to wander freely, but when they are constrained, given roles, forced to leave audit trails, and made answerable to human review.

This section has described how the audition is staged and adjudicated. The next section turns from apparatus to what the audition provisionally reveals: model performance, reliability, role-weighted scoring, and the cautious bridge from benchmark evidence to STA base-model selection.

## What the Audition Reveals

The culmination of MARS-Bench is not a leaderboard. A leaderboard is one of its artifacts, but the audition’s purpose is not to crown the most fluent model; it is to determine which models are plausible substrates for later STA formation, and to do so without confusing raw score, role-fidelity, failure resilience, licensing, architecture, and fine-tuning feasibility. This section, then, reports the complete Codex-Claude statistical audit for the current MARS-Bench release, while preserving the distinction between comparator models, which establish ceilings and architectural contrasts, and STA-candidate models, which are practically eligible for the downstream PhD pipeline.

The replication package archives the full release described above (251 questions, 10,542 responses) along with judging rows, reference and rubric materials, judge notes, audit flags, provenance hashes, and the statistical scripts sufficient to regenerate the final results.<sup>31</sup> Of the 10,542 release responses, 9,618 are paired Codex–Claude rows over the 229 questions analyzed below; the remaining 924 are scored by a dedicated D1 structured-morphology scorer (admitting deterministic adjudication for that question type) and are excluded from the paired statistics unless a morphology-specific bridge is declared.

---

<sup>30</sup> Andrew G. Elrod, *minsky: A Sequential, Multi-CLI Adversarial-Audit Harness for High-Stakes Scholarly Work*, GitHub repository, 2026, concept DOI 10.5281/zenodo.20021747

<sup>31</sup> Andrew G. Elrod, *MARS-Bench Statistical Replication Package*, version 2026.05.07, GitHub/Zenodo, DOI: 10.5281/zenodo.20067936. This public export removes private user/account tables, worker-claim state, machine-local paths, and raw long-form source/context bodies while retaining hashes and structured provenance.

## Judge reliability

Before model scores can be interpreted, the two judge lanes must be shown to be measuring the same response-quality construct. The completed paired dataset supports that claim at the aggregate level. The two-way mixed-effects, absolute-agreement intraclass correlation is high for a single judge,  $ICC(A,1) = 0.9267$ , and higher still for the averaged two-judge score,  $ICC(A,2) = 0.962$ , with a bootstrap 95% confidence interval of 0.9573–0.9660 (1000 percentile-interval replicates, seed 20260503).<sup>32</sup> The mean signed judge difference, Claude minus Codex, is 0.0271 on the 0–5 scale, while the mean absolute difference is 0.2261. Pearson and Spearman correlations are similarly high, at 0.9312 and 0.9218 respectively; 85.39% of paired rows fall within 0.50 points and 98.33% within 1.00 point.




Figure 4: Bland-Altman agreement of the two judge lanes (Codex CLI, gpt-5.5-medium-codex-cli; Claude Code CLI, claude-opus-4-7-medium-claude-cli) across  $n = 9,618$  paired rows. Hex-bin density encodes paired-row counts on a logarithmic scale. The mean signed difference (Claude – Codex) is +0.0271; the parametric 95% limits of agreement are  $[-0.6921, +0.7464]$ ; the empirical 2.5–97.5 percentile band is  $[-0.7, +0.9]$ . The narrow concentration of mass around zero across the score range supports use of the averaged Codex–Claude score for aggregate model

<sup>32</sup> For  $ICC(A,1)$  and  $ICC(A,k)$ , see Patrick E. Shrout and Joseph L. Fleiss, "Intraclass Correlations: Uses in Assessing Rater Reliability," *Psychological Bulletin* 86, no. 2 (1979): 420–28; Kenneth O. McGraw and S. P. Wong, "Forming Inferences About Some Intraclass Correlation Coefficients," *Psychological Methods* 1, no. 1 (1996): 30–46. For the bootstrap procedure used here, see Bradley Efron and Robert J. Tibshirani, *An Introduction to the Bootstrap* (New York: Chapman and Hall, 1993). All bootstrap intervals reported in this paper, including the role-weighted composite confidence intervals in the leaderboard below and the paired Llama-Gemma difference, were computed under the same configuration; the SHA-256 fingerprint of the paired-rows input is f61d100c..., recorded in full in the replication package's cluster\_bootstrap\_results.json metadata.

comparison, while the tail width preserves row-level disagreement as a target for qualitative audit, especially in source-critical domains.

This is strong evidence for using the averaged Codex-Claude score in aggregate model selection. It is not evidence that every row is settled. Bland-Altman limits remain wide enough to matter at the level of individual disputed answers: the normal limits are approximately  $-0.6921$  to  $+0.7464$ , and the empirical limits are  $-0.7$  to  $+0.9$ .<sup>33</sup> That distinction is important. MARS-Bench can justify an averaged score for model comparison while still treating high-disagreement rows, especially in D6 and D7, as audit targets. The audition has reliable adjudicators; it does not thereby abolish the need for the director’s eye.

## Role-audition weighting

Raw means are informative but not sufficient. D1 contains many philological items, while D7 bears the paper’s central Marcion Studies role-pressure; a simple row-count average would let the benchmark’s item distribution determine the model-selection theory. MARS-Bench therefore reports a declared role-audition weighted composite:

$$\text{Composite} = 0.08 \cdot D_1 + 0.17 \cdot D_2 + 0.175 \cdot D_3 + 0.07 \cdot D_4 + 0.14 \cdot D_5 + 0.145 \cdot D_6 + 0.22 \cdot D_7$$

These weights give the most influence to Marcion Studies, patristic knowledge, biblical/canonical reasoning, and heresiological source discipline, while still retaining philology, historical context, and constructive theology. The weights are not metaphysical; they are a declared selection instrument tied to the downstream STA role. Their purpose is to keep the audition from selecting a model merely because it is broadly good at language questions, or broadly verbose in theological prose, while failing at the role’s most dangerous source-critical pressures.

| Rank | Model                                 | Role-weighted composite | 95% CI           | Top-3 bootstrap rate | STA-selection status                     |
|------|---------------------------------------|-------------------------|------------------|----------------------|------------------------------------------|
| 1    | hermes-3-llama-3.1-405b <sup>34</sup> | 2.7995                  | [2.7222, 2.8796] | 1.000                | Comparator; >70B / currently impractical |

<sup>33</sup> J. Martin Bland and Douglas G. Altman, “Statistical Methods for Assessing Agreement Between Two Methods of Clinical Measurement,” *The Lancet* 327, no. 8476 (1986): 307–10.

<sup>34</sup> Ideally, the original llama-3.1-405b model would have been tested. However, at the time of presenting this paper, no viable option for running this model could be found. The project does not currently have the available compute resources to run a dense model of this size locally, and no API service could be located. Only the NousResearch Hermes fine-tuned variant was available (<https://huggingface.co/NousResearch/Hermes-4-405B>) and run through the Deepinfra API.

| Rank | Model                        | Role-weighted composite | 95% CI           | Top-3 bootstrap rate | STA-selection status                                |
|------|------------------------------|-------------------------|------------------|----------------------|-----------------------------------------------------|
| 2    | llama4-maverick-17b-128e-moe | 2.5557                  | [2.4975, 2.6176] | 0.999                | Comparator; MoE                                     |
| 3    | llama-3.3-70b-instruct       | 2.4952                  | [2.4335, 2.5568] | 0.873                | STA candidate; strongest role-fidelity comparator   |
| 4    | gemma-4-31b-it               | 2.4566                  | [2.4048, 2.5051] | 0.128                | STA candidate; preferred practical substrate        |
| 5    | qwen3.5-397b-a17b-moe        | 2.3978                  | [2.3595, 2.4372] | 0.000                | Comparator; large MoE                               |
| 6    | deepseek-v4-flash            | 2.3830                  | [2.3434, 2.4245] | 0.000                | Comparator; future judge-lane interest              |
| 7    | glm-5.1                      | 2.3681                  | [2.3214, 2.4108] | 0.000                | Comparator; large/MoE                               |
| 8    | command-a-03-2025            | 2.3452                  | [2.2916, 2.4060] | 0.000                | Comparator; non-commercial license excludes STA use |
| 9    | qwen3.5-122b-a10b            | 2.2771                  | [2.2333, 2.3206] | 0.000                | Comparator; MoE                                     |
| 10   | qwen3.6-35b-a3b-q8_0         | 2.2626                  | [2.2191, 2.3092] | 0.000                | Comparator; MoE / RI follow-up                      |
| 11   | mistral-medium-3.5           | 2.2370                  | [2.1914, 2.2799] | 0.000                | Comparator; >70B / currently impractical            |
| 12   | qwen3.6-27b-fp8              | 2.2252                  | [2.1899, 2.2653] | 0.000                | STA-eligible dense; not selected                    |
| 13   | llada2.1-flash <sup>35</sup> | 2.0532                  | [2.0173, 2.0936] | 0.000                | Comparator; MoE                                     |
| 14   | qwen3.5-9b                   | 1.9451                  | [1.9110, 1.9811] | 0.000                | STA-eligible dense; not selected                    |

<sup>35</sup> LLaDa 2.1 Flash is a diffusion language model (DLM), not a transformer model (LLM). As such, it is currently not an eligible candidate to serve as an STA substrate. Instead, it is presented here as an architectural comparator. Performance for this model was low, but given future advances in the diffusion architecture space, it may be relevant to consider diffusion models as a realistic option. This is worth exploring, particularly as this architecture has proven to outperform autoregressive transformer models when limited training data is available: cf. Jinjie Ni et al., “Diffusion Language Models Are Super Data Learners,” version 1, preprint, arXiv, 2025, <https://doi.org/10.48550/ARXIV.2511.03276>.

The first result is unsurprising but useful: hermes-3-llama-3.1-405b establishes a clear ceiling. It ranks first on the role-weighted composite, has a 1.0 top-1 and top-3 bootstrap rate, and performs especially well in the source-critical center of the benchmark. But it is not a practical STA base-model candidate in the present project. It is a comparator: a way of seeing how far current large-scale open-weight systems can go when compute and architecture constraints are relaxed.

The second result is more interesting. llama4-maverick-17b-128e-moe ranks second and has a near-certain top-3 bootstrap rate, but it remains a MoE comparator rather than the immediate downstream choice. Its observed advantage over the practical dense candidates is real in the descriptive leaderboard, but the margin is not large enough, by itself, to override the selection bridge. The paper should not claim that larger or MoE systems are irrelevant. But the present statistics do not make the case that the downstream STA project should abandon its dense, fine-tunable, open-science-oriented constraints merely to chase a modest comparator gain.

## The eligible-candidate decision: Llama 3.3 versus Gemma 4

Once the comparator cohort is separated from the STA-candidate cohort, the practical decision narrows. llama-3.3-70b-instruct is the strongest role-weighted STA candidate, with a composite of 2.4952 and a bootstrap top-3 rate of 0.873. gemma-4-31b-it follows closely, with a composite of 2.4566 and a top-3 rate of 0.128. The observed difference is 0.0386 in Llama's favor. The targeted paired bootstrap test, however, gives a 95% confidence interval of  $-0.0251$  to  $+0.1052$ , which includes zero.<sup>36</sup> On this snapshot, then, Llama's role-weighted advantage over Gemma is not statistically decisive under the declared paired-difference procedure.

The qualitative direction is visible in Figure 4: Llama 3.3 leads in the role-pressured domains (D3 patristic, D4 early Christian history, D6 heresiological, D7 Marcion Studies), while Gemma 4 leads in the formation-substrate domains (D1 ancient languages, D2 biblical, D5 theological). That finding should be used carefully. It does not mean the two models are identical. Llama is stronger in the most role-specific regions of the benchmark: D7 Marcion Studies, D6 heresiological discipline, D4 early Christian history, and D3 patristic knowledge. Gemma is stronger in D1 philology, D2 biblical/canonical substrate, and D5 theological reasoning. The difference is therefore qualitative as well as numerical. Llama looks more immediately role-pressured;

---

<sup>36</sup> cf. Nathaniel Schenker and Jane F. Gentleman, "On Judging the Significance of Differences by Examining the Overlap Between Confidence Intervals," *The American Statistician* 55, no. 3 (2001): 182–86. For an equivalence-testing approach (TOST) that converts a "CI includes zero" reasoning into a positive statement of practical equivalence, see "Between Confidence Intervals," *The American Statistician* 55, no. 3 (2001): 182–86. For an equivalence-testing approach (TOST) that converts a "CI includes zero" reasoning into a positive statement of practical equivalence, see Daniël Lakens, "Equivalence Tests: A Practical Primer for t Tests, Correlations, and Meta-Analyses," *Social Psychological and Personality Science* 8, no. 4 (2017): 355–62.

Gemma looks more attractive as a compact, dense, permissively licensed, theologically and philologically competent substrate. Because the score difference is not decisive, the non-score criteria legitimately become part of the decision rather than post-hoc excuses for ignoring the leaderboard.




Figure 5: Mean averaged Codex-Claude score by model and domain (n = 9,618 paired rows). Rows are ordered by the role-weighted composite (descending), matching the leaderboard above. Cell values are the mean of the per-row averaged judge scores within each model-domain pair, on the rubric's 0-5 scale. The Llama 3.3-Gemma 4 asymmetry that the prose argues qualitatively is visible directly: Gemma 4 leads in D1, D2, and D5; Llama 3.3 leads in D3, D4, D6, and D7. Hermes-3-Llama-3.1-405B (top row) sets the ceiling across most domains; the bottom rows (LLaDA 2.1 and Qwen3.5 9B) display the systematic shortfall expected of much smaller or non-autoregressive systems.

Those criteria matter for this PhD project. The downstream STA architecture requires weight access for fine-tuning, preferably under open-science-compatible licensing; it also benefits from a dense model small enough for LoRA supervised fine-tuning, possible later continued pretraining, and experimental hidden-state or logit interventions. Under that bridge, gemma-4-31b-it remains the defensible provisional base-model selection: not because it wins the benchmark, but because it remains close to Llama 3.3 on the role-weighted composite while offering a more favorable license, smaller dense architecture, lower compute burden, clean derivative naming, and zero observed response-integrity failures. Llama 3.3 remains the nearest role-

fidelity comparator and should remain in the experimental orbit, especially if licensing or compute conditions change.

## Failure profile, stability, and verbosity

The failure diagnostics reinforce the point that model selection cannot be reduced to a score column. The twelve response-integrity failures profiled in §Response-integrity gates above are rare but not absent, and they are concentrated unevenly: four of the five affected models are MoE or quantized, and five of the twelve failures fall in D7, the most role-pressured domain. Models with zero observed failures cannot be inferred immune, but they pass the diagnostic and are eligible for the next stage of the audition.<sup>37</sup>

Run-to-run variance supplies another caution. A model can be impressive and unstable. hermes-3-llama-3.1-405b, despite leading the benchmark, has the highest mean coefficient of variation in the run-to-run diagnostic, while Gemma's mean CV is lower and among the more stable profiles. Llama 3.3 is less stable than Gemma on this diagnostic. A surrogate who performs brilliantly on some runs and falls out of the role on others is costly to rehearse, difficult to evaluate, and risky in an interactive scholarly setting.

The length diagnostic is also salutary. Long theological prose can sound learned even when it is merely expansive, so a humanities benchmark must test whether verbosity is being rewarded. The length-adjustment model reports a negative log-word coefficient,  $-0.450366$ , with  $p < 0.001$ :  $p \approx 1.3 \times 10^{-297}$ , after controlling for domain, level, and response regime.<sup>38</sup> In other words, the completed audit does not show that longer answers were generally rewarded after those covariates. If anything, the diagnostic suggests that excessive length may correlate with lower scores once difficulty and domain are considered. That does not eliminate the need for human review of fluent pseudo-scholarship, but it weakens one obvious objection to the automated judging apparatus.

## What the audition has shown

Three provisional conclusions follow from the complete Codex-Claude audit. First, the judge apparatus is reliable enough at the aggregate level to support averaged-score

---

<sup>37</sup> Cf. James A. Hanley and Abby Lippman-Hand, "If Nothing Goes Wrong, Is Everything All Right? Interpreting Zero Numerators," *JAMA* 249, no. 13 (1983): 1743–45.

<sup>38</sup> The exact p-value is reported because the diagnostic computes it directly from the t-statistic; " $p < .001$ " would understate the strength of the evidence by ~290 orders of magnitude. Standard caveats about p-values at this scale apply: with  $N = 9,618$  paired rows, even small structural patterns can produce extreme p-values, and the substantive finding here is the sign and magnitude of the coefficient, not a hypothesis-testing decision.

model comparison, while preserving row-level disagreement as a target for qualitative audit. Second, the comparator cohort matters: larger and MoE models establish real ceilings and architectural contrasts, especially Hermes and Llama 4 Maverick. Third, the practical STA-base decision remains a multi-criteria bridge. MARS-Bench does not simply select the highest-scoring system. It selects among models that can be cast.

As such, the result is deliberately double. If the question is “which model performed best in this audition?” the answer is hermes-3-llama-3.1-405b. If the question is “which currently practical substrate should be carried forward into STA rehearsal under the project’s open-science, dense-architecture, and compute constraints?” the answer remains provisionally gemma-4-31b-it, with llama-3.3-70b-instruct retained as the strongest eligible role-fidelity comparator. This is the benchmark doing what an audition must do: distinguishing virtuosity from castability.

The next section turns from findings to validity and limits. MARS-Bench has shown that an agentic, source-grounded benchmark can produce auditable and statistically coherent evidence for humanities model selection. It has not yet shown that the selected base model, once fine-tuned, will produce successful STAs. That stricter predictive claim belongs to the next phase of the PhD project, where MARS-Bench scores must be tested against expert-rated STA performance.

## Validity, Limits, and the Next Rehearsal

Validity is not a magic property conferred by a benchmark score. It is the relation between a claim and the evidence that claim can bear. Hence, MARS-Bench is valid at this stage if it can support a castability claim: which model is worth carrying into STA rehearsal under declared scholarly, statistical, licensing, and compute constraints. It is not valid, and should not be asked to be valid, as proof that the selected model can already surrogate Marcionite theological reasoning in the full sense required by the PhD project. The audition can select an actor; it cannot complete the performance.

The benchmark’s validity begins in its refusal to treat Marcionite surrogation as a general knowledge task. D1–D7 do not measure generic intelligence, generic theology, or general patristic fluency. They test the role-world: ancient languages, scriptural substrate, patristic control, early Christian historical scene-setting, theological discrimination, heresiological source discipline, and direct Marcion Studies pressure. The rubric’s movement from factual recall through inferential paraphrase to synthetic reasoning likewise constrains the claim. Recall and paraphrase test formation; synthesis tests whether that formation can be recomposed under pressure. D6/D7 witness discipline, response-integrity gates, and trap caps then prevent fluent impersonation from being mistaken for role-eligible performance. MARS-Bench is strongest when understood in precisely this narrower way: as a source-critical



audition for constrained theological generation, not as a general-purpose theological intelligence test.<sup>39</sup>

This selection bridge also has declared boundaries. The paired Codex–Claude model-selection analysis covers 229 questions and 9,618 paired response rows. The full release contains 251 questions and 10,542 responses, but 22 D1 structured morphology adjunct questions were scored by a dedicated morphology scorer rather than by the paired Codex–Claude lanes, and are therefore outside the paired model-selection statistics. Current-version post-benchmark expert spot-checking also remains the next audit phase. Earlier human spot checks materially improved the benchmark, especially in D6 and D7, but they should not be represented as a completed expert validation of the final scored version. These qualifications define the scope within which the method should be trusted.

Within that scope, Gemma 4 31B-IT is the provisional practical STA base model. This is not because Gemma *wins* MARS-Bench. The role-weighted composite is led by large comparator models, especially hermes-3-llama-3.1-405b and llama4-maverick-17b-128e-moe, whose results are valuable as ceiling and architecture comparators but not presently actionable under the project’s stated constraints. Among the currently practical dense candidates, Llama 3.3 has the observed role-weighted advantage over Gemma 4, but the paired bootstrap difference is small and not statistically decisive (the full  $\Delta$  and 95% CI are reported in §The eligible-candidate decision above).

That choice is practical, legal, and methodological. Gemma’s dense 31B profile reduces the compute burden of fine-tuning and repeated STA experimentation when compared with a 70B model, while still remaining within the performance neighborhood of Llama 3.3 on the role-weighted composite. Its run-to-run behavior is also somewhat more stable in the current diagnostics: Gemma records lower mean CV (0.0860 versus Llama’s 0.1055), lower mean variance (0.11006 versus 0.14590), and fewer high-CV questions above the 0.15 threshold (35 versus 50). Both models record zero observed response-integrity hard failures across 753 release responses, which is favorable but not proof of immunity; by the rule-of-three heuristic, the approximate one-sided 95% upper bound is 0.003984.

Licensing matters as well. Gemma 4 is available under Apache 2.0, and Google’s model documentation identifies the 31B dense variant as designed to bridge server-grade performance and local execution.<sup>40</sup> Llama 3.3 70B Instruct, by contrast, is distributed

---

<sup>39</sup> For the public replication package, see Andrew G. Elrod, *MARS-Bench Statistical Replication Package*, version 2026.05.07, GitHub/Zenodo, DOI: 10.5281/zenodo.20067936.

<sup>40</sup> Google Developers Blog, “Bring state-of-the-art agentic skills to the edge with Gemma 4,” 2 April 2026; Google AI for Developers, “Gemma 4 model overview”; Hugging Face, google/gemma-4-31B-it model card. The Hugging Face card lists license apache-2.0; Google’s overview describes the 31B model as a dense model bridging server-grade performance and local execution.

under the custom Llama 3.3 Community License Agreement.<sup>41</sup> Nothing in this paper claims that VU, the EU, or open-science norms require Apache 2.0 model weights. The claim is narrower: a permissive, familiar license better serves the project’s reproducibility, derivative-release, and dissemination aims. The VU’s Research Data and Software Management Policy foregrounds visibility, verifiability, reusability, FAIR management, persistent identifiers, metadata, and licensing for research data and software; the VU’s research-software guidance likewise treats version control, documentation, and licensing as minimum quality expectations; and Horizon Europe frames open science around open access and FAIR data under the principle “as open as possible, as closed as necessary.”<sup>42</sup> In the performance parlance of this paper, Gemma is an actor who can legally take the contract, fit the production’s budget, rehearse reliably, and appear on stage under the production’s actual conditions. It is logistically and virtuosically suited for the role.

## Audition versus Actor

The casting decision is made; rehearsal is next. The next phase is rehearsal: fine-tuning, corpus-memory, retrieval infrastructure where appropriate, persona-specific workflows, and expert-directed evaluation. The aforementioned four experimental STA personae should be understood as different stagings of the same surrogate problem. Each persona asks the selected substrate to perform under different given circumstances: exegetical, dialogical, historical-counterfactual, and constructive-theological.

Figure 6 transposes Stanislavski's system into this downstream STA architecture: MARS-Bench functions as casting and base-model selection; fine-tuning, provenance, and corpus memory form the rehearsal apparatus; persona-specific workflows stage the expert encounter; and the proposed supertask is to surrogate Marcion's absent voice under scholarly constraint.

---

<sup>41</sup> Meta, “Llama 3.3 Community License Agreement.” The model card lists license llama3.3 and describes the custom Llama 3.3 Community License Agreement.

<sup>42</sup> Vrije Universiteit Amsterdam, *Research Data and Software Management Policy*, version 3.0, 17 December 2024; Vrije Universiteit Amsterdam, “Meet the Research Software quality standards”; VU Research Support Handbook, “Software Licensing”; European Commission, “Open science,” Horizon Europe guidance. The policy/guidance is used here to support alignment with reproducibility, metadata, licensing, and FAIR/open-science norms, not to claim that Apache 2.0 is legally mandatory for model weights.




Figure 6: The downstream STA architecture as a Stanislavskian rehearsal apparatus. MARS-Bench occupies the casting / base-model-selection step (left); the rehearsal apparatus comprises fine-tuning, provenance discipline, and corpus memory (center); persona-specific workflows stage the expert encounter under different given circumstances (right). The figure describes the next phase of the PhD project, not a result already demonstrated by MARS-Bench.

operationally stageable.

## Exeunt, Manet Marcion

To determine a student's competency, an assessment must be designed around the particulars of the subject matter. To cast an actor, the audition must live within the constraints of the role. To select an LLM to serve as the substrate for an esoteric, nuanced objective, only a bespoke benchmark will do. None of these procedures, however, pretends to recover what is already constituted in the candidate. The student is not invented by the examination. The actor is not summoned from the dead by the audition. The model is not Marcion. What each procedure accomplishes is narrower, and for that reason more honest: it nominates eligibility, under declared conditions, for a role that remains to be performed.

The remaining limits, then, are best understood as next rehearsal conditions. Predictive validity remains untested: a base-model audition does not guarantee fine-tuned STA success. Source-packet selectivity still matters because source-grounded judging is only as strong as the source-book it uses. Automated judges may share blind spots even when inter-judge reliability is high. Pretraining exposure remains double-edged: it is formation for the role, but also a possible source of benchmark contamination. Roach and Stanislavski warrant the method's grammar of substitution, given circumstances, and audition; they do not warrant ontological identity with the lost historical actor.

MARS-Bench, therefore, reaches a disciplined stopping point. It validates a selection procedure, not a resurrected voice; it supports Gemma as a provisional STA substrate, not as a finished surrogate; and it turns the unresolved problems of expert evaluation, fine-tuned behavior, source-packet selectivity, and persona-specific performance into the next rehearsal conditions. The benchmark is not a step in undoing Marcion's *as was* absence. It has made that absence

Marcion *as was* is not coming back. He has not come back through Tertullian's invective, nor through Epiphanius's catalog, nor through Adamantius's dialogue, nor through Harnack's century-old reconstruction, nor through the patient text-critical labors of Klinghardt, BeDuhn, Roth, Lieu, or Vinzent. He will not come back through a transformer either. The audition does not contest the fact of his absence. It contests only the resignation that absence ordinarily licenses: the assumption that what cannot be recovered cannot be staged, and that what cannot be staged cannot be made newly arguable. The performance theorist, asked to choose between actuality and falsehood, refuses both. Stanislavski's *Director* puts it with arresting plainness: in ordinary life, truth is what really exists, what a person really knows; on the stage it consists of something that is not actually in existence but which could happen.<sup>43</sup> Between the actual and the merely fictive there is a third register, and disciplined performance does its work there.

Roach, writing in a different idiom and toward different ends, arrives at a structurally identical insight. The kinesthetic imagination, he writes, inhabits the realm of the virtual; its truth is the truth of simulation, of fantasy, or of daydreams, but its effect on human action may have material consequences of the most tangible sort and of the widest scope.<sup>44</sup> The pairing is not coincidental. Both Stanislavski and Roach have refused the binary that has shadowed historical-theological work on Marcion since Harnack: either the lost figure can be recovered, or the project is a failure. There is a third option, and it is older than either theorist names. It is the option of *as if*: the rule-bound imaginative condition that authorizes action inside a world one did not make. Scenic truth has scholarly consequences. That is the wager of the audition.

What MARS-Bench has produced, then, is not Marcion, and not a model of Marcion, but a castable substrate: a Gemma 4 31B-IT eligible, under declared statistical, licensing, and architectural constraints, for the rehearsal that follows. What remains is rehearsal proper: fine-tuning, corpus-memory, retrieval where appropriate, the four experimental personae, and above all, the expert direction for which no automated judge can definitively substitute. The benchmark has held open a place under disciplined conditions. A candidate has stepped into it. One of them, provisionally, has been judged eligible to remain. The next task is staging, and the staging will succeed or fail under conditions the audition cannot anticipate.

This is what the introduction promised and what the body of the paper has tried to demonstrate. Ray Charles tested Jamie Foxx not just as a piano player but as a man, because the role required more than fluency in its surface vocabulary. MARS-Bench tests models not as reciters of Marcionite topoi but as candidates for cross-domain performance under hostile-source pressure, because the role of theological surrogate requires more than fluency in heresiological labels. In each case, the audition does not

---

<sup>43</sup> Stanislavski, 141.

<sup>44</sup> Roach, 27.

predict the performance. It warrants the candidate. The performance itself, whether Foxx's Charles or the STAs' Marcion, will be judged elsewhere, by other criteria, under direction the audition cannot provide.

The benchmark has not undone Marcion's absence. It has refused to be governed by it. *As if*, not *as was*: the difference, far from being a concession, is the discipline that makes the work possible at all.

# Bibliography

- Adamantius. *Dialogue on the True Faith in God = De Recta in Deum Fide*. Translated by Robert A. Pretty. Edited by Garry W. Trompf. Gnostica 1. Leuven: Peeters, 1997.
- BeDuhn, Jason. *The First New Testament: Marcion's Scriptural Canon*. Salem, OR: Polebridge Press, 2013.
- Bingham, Dennis. *Whose Lives Are They Anyway? The Biopic as Contemporary Film Genre*. New Brunswick, NJ: Rutgers University Press, 2010.
- Bland, J. Martin, and Douglas G. Altman. "Statistical Methods for Assessing Agreement between Two Methods of Clinical Measurement." *The Lancet* 327, no. 8476 (1986): 307–10.
- Brown, Tom, and Belén Vidal, eds. *The Biopic in Contemporary Film Culture*. New York: Routledge, 2014.
- Carnicke, Sharon Marie. *Stanislavsky in Focus: An Acting Master for the Twenty-First Century*. 2nd ed. London and New York: Routledge, 2009.
- CenterBLC. *CenterBLC/LXX: Rahlfs' LXX 1935 v1.0.0*. Zenodo, 2021. <https://doi.org/10.5281/zenodo.5807148>.
- Chiang, Wei-Lin, Lianmin Zheng, Ying Sheng, Anastasios N. Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E. Gonzalez, and Ion Stoica. "Chatbot Arena: An Open Platform for Evaluating LLMs by Human Preference." arXiv:2403.04132, 2024.
- Clark, Peter, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. "Think You Have Solved Question Answering? Try ARC, the AI2 Reasoning Challenge." arXiv:1803.05457, 2018.
- Doyle, Colin. "LLMs as Method Actors: A Model for Prompt Engineering and Architecture." arXiv:2411.05778, 2024. <https://doi.org/10.48550/arXiv.2411.05778>.
- Efron, Bradley, and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Monographs on Statistics and Applied Probability 57. New York: Chapman and Hall, 1993.
- Elrod, Andrew G. *MARS-Bench Statistical Replication Package*. Version 2026.05.07. GitHub and Zenodo, 2026. <https://doi.org/10.5281/zenodo.20067936>.
- Elrod, Andrew G. *minsky: A Sequential, Multi-CLI Adversarial-Audit Harness for High-Stakes Scholarly Work*. GitHub and Zenodo, 2026. Concept DOI <https://doi.org/10.5281/zenodo.20021747>.

- Epiphanius of Salamis. *The Panarion of Epiphanius of Salamis, Book I (Sects 1–46)*. Translated by Frank Williams. Nag Hammadi and Manichaean Studies 35. Leiden: Brill, 1987.
- Epiphanius of Salamis. *The Panarion of Epiphanius of Salamis, Books II and III, De Fide (Sects 47–80)*. Translated by Frank Williams. Nag Hammadi and Manichaean Studies 36. Leiden: Brill, 1994.
- ETCBC (Eep Talstra Centre for Bible and Computer). *BHSA: Hebrew Bible + Linguistic Annotations in Text-Fabric Format*. DANS Data Station, 2017–.  
<https://doi.org/10.17026/dans-z6y-skyh>.
- European Commission. "Open Science." Horizon Europe Programme Guide, Research and Innovation, n.d. [https://research-and-innovation.ec.europa.eu/strategy/strategy-research-and-innovation/our-digital-future/open-science\\_en](https://research-and-innovation.ec.europa.eu/strategy/strategy-research-and-innovation/our-digital-future/open-science_en).
- Goodale, Gloria. "An Audition for 'Ray' with Ray." *Christian Science Monitor*, October 29, 2004.
- Google. "Bring State-of-the-Art Agentic Skills to the Edge with Gemma 4." Google Developers Blog, April 2, 2026.
- Google. "Gemma 4 Model Overview." Google AI for Developers, n.d.
- Google. "google/gemma-4-31B-it." Hugging Face Model Card, n.d.
- Gu, Jiawei, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, Honghao Liu, et al. "A Survey on LLM-as-a-Judge." arXiv:2411.15594, 2024.
- Hanley, James A., and Abby Lippman-Hand. "If Nothing Goes Wrong, Is Everything All Right? Interpreting Zero Numerators." *JAMA* 249, no. 13 (1983): 1743–45.
- Harnack, Adolf von. *Marcion: Das Evangelium vom fremden Gott. Eine Monographie zur Geschichte der Grundlegung der katholischen Kirche*. Texte und Untersuchungen 45. Leipzig: J. C. Hinrichs, 1921. Reprinted Darmstadt: Wissenschaftliche Buchgesellschaft, 1985.
- Harnack, Adolf von. *Marcion: The Gospel of the Alien God*. Translated by John E. Steely and Lyle D. Bierma. Durham, NC: Labyrinth Press, 1990.
- Hippolytus of Rome. *Refutation of All Heresies*. Translated by M. David Litwa. Writings from the Greco-Roman World 40. Atlanta: SBL Press, 2016.
- Hoffmann, R. Joseph. *Marcion, On the Restitution of Christianity: An Essay on the Development of Radical Paulinist Theology in the Second Century*. American Academy of Religion Academy Series 46. Chico, CA: Scholars Press, 1984.

- Irenaeus of Lyons. *Against the Heresies, Book I*. Translated by Dominic J. Unger, with further revisions by John J. Dillon. Ancient Christian Writers 55. New York: Newman Press, 1992.
- Jurg, Tony, Saulo de Oliveira Cantanhêde, and Oliver Glanz. *CenterBLC/N1904: 1.0.0*. Zenodo, 2024. <https://doi.org/10.5281/zenodo.13117911>.
- Justin Martyr. *Justin, Philosopher and Martyr: Apologies*. Edited and translated by Denis Minns and Paul Parvis. Oxford Early Christian Texts. Oxford: Oxford University Press, 2009.
- Kemp, Richard J. *Embodied Acting: What Neuroscience Tells Us about Performance*. London and New York: Routledge, 2012.
- Kim, Seungone, Juyoung Suk, Shayne Longpre, Bill Yuchen Lin, Jamin Shin, Sean Welleck, Graham Neubig, Moontae Lee, Kyungjae Lee, and Minjoon Seo. "Prometheus 2: An Open Source Language Model Specialized in Evaluating Other Language Models." In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. arXiv:2405.01535, 2024.
- Kinzig, Wolfram. *Harnack, Marcion und das Judentum: Nebst einer kommentierten Edition des Briefwechsels Adolf von Harnacks mit Houston Stewart Chamberlain*. Arbeiten zur Kirchen- und Theologiegeschichte 13. Leipzig: Evangelische Verlagsanstalt, 2004.
- Klinghardt, Matthias. *Das älteste Evangelium und die Entstehung der kanonischen Evangelien: Untersuchung – Rekonstruktion – Übersetzung – Varianten*. 2 vols. Texte und Arbeiten zum neutestamentlichen Zeitalter 60. Tübingen: Francke, 2015.
- Klinghardt, Matthias. *The Oldest Gospel and the Formation of the Canonical Gospels: Inquiry – Reconstruction – Translation – Variants*. Translated by the author. 2 vols. Biblical Tools and Studies 41–42. Leuven: Peeters, 2021.
- Krumdick, Michael, Charles Lovering, Varshini Reddy, Seth Ebner, and Chris Tanner. "No Free Labels: Limitations of LLM-as-a-Judge without Human Grounding." arXiv:2503.05061, 2025.
- Lakens, Daniël. "Equivalence Tests: A Practical Primer for *t* Tests, Correlations, and Meta-Analyses." *Social Psychological and Personality Science* 8, no. 4 (2017): 355–62.
- Lieu, Judith M. *Marcion and the Making of a Heretic: God and Scripture in the Second Century*. Cambridge: Cambridge University Press, 2015.
- Lin, Stephanie, Jacob Hilton, and Owain Evans. "TruthfulQA: Measuring How Models Mimic Human Falsehoods." In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 3214–52. Dublin: Association for Computational Linguistics, 2022.



- Litwa, M. David. *The Evil Creator: Origins of an Early Christian Idea*. New York: Oxford University Press, 2021.
- Litwa, M. David. *Found Christianities: Remaking the World of the Second Century CE*. London: Bloomsbury T&T Clark, 2022.
- McGraw, Kenneth O., and S. P. Wong. "Forming Inferences about Some Intraclass Correlation Coefficients." *Psychological Methods* 1, no. 1 (1996): 30–46.
- Meta. "Llama 3.3 Community License Agreement." Meta AI, n.d. [https://www.llama.com/llama3\\_3/license/](https://www.llama.com/llama3_3/license/).
- Moll, Sebastian. *The Arch-Heretic Marcion*. Wissenschaftliche Untersuchungen zum Neuen Testament 250. Tübingen: Mohr Siebeck, 2010.
- Nederlandse Organisatie voor Wetenschappelijk Onderzoek (NWO). "Interpreting the Unspoken: Counterfactual Historical Reasoning for Reconstructing Early Christian Heterodox Thought Using Large Language Models." Doctoral Grant for Teachers (Promotiebeurs voor Leraren), n.d. <https://doi.org/10.61686/DQUQD55640>.
- Ngũgĩ wa Thiong'o. *Decolonising the Mind: The Politics of Language in African Literature*. London: James Currey, 1986.
- Ni, Jinjie, Qian Liu, Longxu Dou, Chao Du, Zili Wang, Tianyu Pang, Wenjun Lou, et al. "Diffusion Language Models Are Super Data Learners." Version 1. Preprint. arXiv:2511.03276, 2025. <https://doi.org/10.48550/arXiv.2511.03276>.
- Nous Research. "NousResearch/Hermes-4-405B." Hugging Face Model Card, n.d. <https://huggingface.co/NousResearch/Hermes-4-405B>.
- NPR. "Playing a 'Genius': Jamie Foxx Takes on Ray Charles." *NPR*, October 22, 2004.
- Ogunnaike, Lola. "Jamie Foxx Goes Dark to Play a Musical Hero." *New York Times*, September 12, 2004, sec. 2, p. 33.
- OpenCode. *OpenCode: Agentic Coding Harness*. GitHub repository, n.d. <https://github.com/anomalyco/opencode>.
- Panickssery, Arjun, Samuel R. Bowman, and Shi Feng. "LLM Evaluators Recognize and Favor Their Own Generations." In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*. arXiv:2404.13076, 2024.
- Phan, Long, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, Hugh Zhang, Chen Bo Calvin Zhang, Mohamed Shaaban, John Ling, Sean Shi, et al. "Humanity's Last Exam." arXiv:2501.14249, 2025.

- Rein, David, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. "GPQA: A Graduate-Level Google-Proof Q&A Benchmark." arXiv:2311.12022, 2023.
- Roach, Joseph. *Cities of the Dead: Circum-Atlantic Performance*. New York: Columbia University Press, 1996.
- Roorda, Dirk. *Text-Fabric*. Software, 2016–. <https://doi.org/10.5281/zenodo.592193>.
- Roorda, Dirk. "Text-Fabric: Handling Biblical Data with IKEA Logistics." *HIPHIL Novum* 5, no. 2 (2019): 126–35. <https://doi.org/10.7146/hn.v5i2.142740>.
- Roth, Dieter T. *The Text of Marcion's Gospel*. New Testament Tools, Studies, and Documents 49. Leiden: Brill, 2015.
- Schenker, Nathaniel, and Jane F. Gentleman. "On Judging the Significance of Differences by Examining the Overlap between Confidence Intervals." *The American Statistician* 55, no. 3 (2001): 182–86.
- Schmid, Ulrich. *Marcion und sein Apostolos: Rekonstruktion und historische Einordnung der marcionitischen Paulusbriefausgabe*. Arbeiten zur neutestamentlichen Textforschung 25. Berlin: De Gruyter, 1995.
- Shrout, Patrick E., and Joseph L. Fleiss. "Intraclass Correlations: Uses in Assessing Rater Reliability." *Psychological Bulletin* 86, no. 2 (1979): 420–28.
- Stanislavski, Constantin. *An Actor Prepares*. Translated by Elizabeth Reynolds Hapgood. London and New York: Routledge, 1989. Originally published New York: Theatre Arts, 1936.
- Stanislavskii, K. S. *Sobranie sochinenii v vos'mi tomakh*. Vol. 3, *Rabota aktera nad soboi v tvorchestvom protsesse perezhivaniia*. *Dnevnik uchenika*. Moscow: Iskusstvo, 1955.
- Tertullian. *Adversus Marcionem, Books I–V*. Edited and translated by Ernest Evans. 2 vols. Oxford Early Christian Texts. Oxford: Clarendon Press, 1972.
- Verga, Pat, Sebastian Hofstatter, Sophia Althammer, Yixuan Su, Aleksandra Piktus, Arkady Arkhangorodsky, Minjie Xu, Naomi White, and Patrick Lewis. "Replacing Judges with Juries: Evaluating LLM Generations with a Panel of Diverse Models." arXiv:2404.18796, 2024.
- Vinzent, Markus. *Marcion and the Dating of the Synoptic Gospels*. Studia Patristica Supplements 2. Leuven: Peeters, 2014.
- Vrije Universiteit Amsterdam. "Meet the Research Software Quality Standards." VU Research Support Portal, n.d.

- Vrije Universiteit Amsterdam. *Research Data and Software Management Policy*. Version 3.0. December 17, 2024.
- Vrije Universiteit Amsterdam. "Software Licensing." VU Research Support Handbook, n.d.
- Wang, Yubo, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, et al. "MMLU-Pro: A More Robust and Challenging Multi-Task Language Understanding Benchmark." In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*. arXiv:2406.01574, 2024.
- Wataoka, Koki, Tsubasa Takahashi, and Ryokan Ri. "Self-Preference Bias in LLM-as-a-Judge." Paper presented at the NeurIPS 2024 Safe Generative AI Workshop, Vancouver, December 2024. arXiv:2410.21819.
- Whorf, Benjamin Lee. "Science and Linguistics." *Technology Review* 42, no. 6 (1940): 229–31, 247–48. Reprinted in *Language, Thought, and Reality: Selected Writings of Benjamin Lee Whorf*, edited by John B. Carroll, 207–19. Cambridge, MA: MIT Press, 1956.
- Whyman, Rose. *The Stanislavski System of Acting: Legacy and Influence in Modern Performance*. Cambridge: Cambridge University Press, 2008.
- Ye, Seonghyeon, Doyoung Kim, Sungdong Kim, Hyeonbin Hwang, Seungone Kim, Yongrae Jo, James Thorne, Juho Kim, and Minjoon Seo. "FLASK: Fine-Grained Language Model Evaluation Based on Alignment Skill Sets." arXiv:2307.10928, 2023.
- You, Runyang, Hongru Cai, Caiqi Zhang, Qiancheng Xu, Meng Liu, Tiezheng Yu, Yongqi Li, and Wenjie Li. "Agent-as-a-Judge." arXiv:2601.05111, 2026.
- Yu, Fangyi. "When AIs Judge AIs: The Rise of Agent-as-a-Judge Evaluation for LLMs." arXiv:2508.02994, 2025.
- Zheng, Lianmin, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. "Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena." In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*. arXiv:2306.05685, 2023.