
AS IF, *not as was*

Auditioning Large Language Models to Surrogate Marcion of Sinope

MARS-Bench · a source-grounded, agentic audition

Andrew Elrod

Vrije Universiteit Amsterdam · HZ University of Applied Sciences

a.g.elrod@vu.nl

May 12, 2026

Today's Program

Five short acts, one intermission of questions

Act I

The Problem

Why a benchmark for an absent heretic?

Act II

The Audition

Stanislavski, Roach, and Jamie Foxx walk into a studio

Act III

The Apparatus

MARS-Bench: domains, packets, agentic judges

Act IV

The Verdict

What 9,618 paired judgments reveal

Act V

Curtain & next steps

Limits, validity, and what the paper does that this talk can't

A C T I

The Problem

“Marcion of Sinope has been silent for eighteen centuries.”

Not absence. Erasure.

The Marcionite archive that was deliberately not kept

What we do not have

- Marcion's own writings — including the Antitheses.
- His edited Gospel and Apostolikon, in any direct manuscript.
- Any Marcionite-internal liturgy, commentary, or correspondence.
- Any sympathetic biographical record of the man himself.

What we do have

Tertullian's five-volume rebuttal. Epiphanius's heresy catalog. Irenaeus, Hippolytus, Adamantius, Theodoret, John of Damascus.

i.e., his enemies — quoting him in order to refute him.

“What Pontic mouse ever had such gnawing powers as he who has gnawed the Gospels to pieces? Verily, O Euxine, you have produced a monster...”

— Tertullian, Adv. Marcionem

Erasure is methodologically different from absence.
Absence is what the archaeologist hopes to fill. Erasure is the absence that was produced — copied, catalogued, and then not preserved. The work of recovery is correspondingly different.

A century of reconstruction

Marcion Studies as the management of a deletion

1921

Adolf von Harnack

Marcion: Das Evangelium vom fremden Gott. Sets the modern reconstructive agenda; treats Marcion as a serious theological figure rather than a heresiological caricature.

1984–2010

Hoffmann · Schmid · Kinzig · Moll

Postwar text-critical and historical refinement. Reconstructions diverge. The figure begins to multiply.

2013–2024

BeDuhn · Klinghardt · Roth · Lieu · Vinzent

Contested arguments about whether Marcion's Gospel is a mutilated Luke, a pre-Lukan source, or something else entirely. The scholarly Marcions no longer agree on basic facts.

The result, after a century: not one Marcion but a contested reconstruction space populated by several plausible Marcions. This is what an LLM has to enter, not avoid.

Interpreting the Unspoken

The PhD project this paper belongs to

*The question: Can LLMs, trained and constrained by carefully curated primary and secondary corpora, be developed into **Synthetic Theological Agents** — synthetic actors capable of surrogating Marcionite theological reasoning under specified historical, philological, and theological circumstances?*

An STA is not

- A recovered Marcionite.
- A machine possessed by the past.
- A claim to have undone the erasure.
- A model fine-tuned to “sound like” Marcion.

An STA is

- A synthetic actor surrogating Marcionite reasoning.
- Operating inside specified historical, philological, theological constraints.
- Inspectable, corrigible, source-answerable.
- Designed to make Marcionite reasoning newly arguable — not to settle it.

Where MARS-Bench sits

Today's talk is the first stage of a four-stage pipeline



In Stanislavski's vocabulary: Stage 1 is casting. Stage 2 is rehearsal. Stage 3 is staging. Stage 4 is the expert direction no automated judge can substitute for. The supertask: surrogate Marcion's absent voice under scholarly constraint.

Today is Stage 1. Stages 2–4 are the next eighteen months of the PhD.

A C T I I

The Audition

“The kid’s got it.” — Ray Charles, allegedly

Stanislavski: constraint as freedom

Live the role inside its “given circumstances”

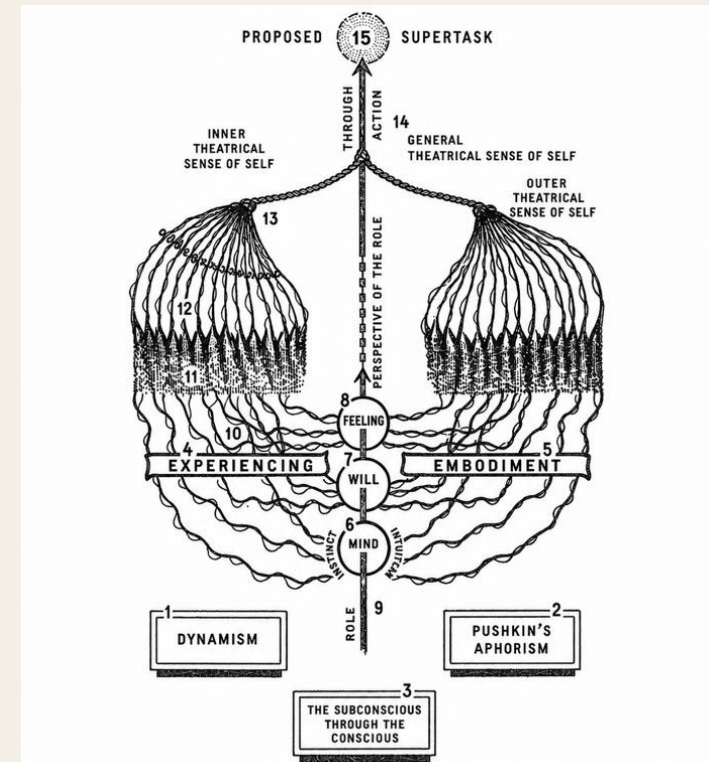
An actor doesn't impersonate. They live the role inside its given circumstances — the story, the social world, the script, the director's reading.

The “if”

Not free-floating fantasy. A rule-bound imaginative condition that authorizes action inside a world one did not make.

Translated into MARS-Bench:

I'm not asking the model to be Marcion. I'm asking whether it can reason as if within the historically mediated circumstances of the Marcion problem.

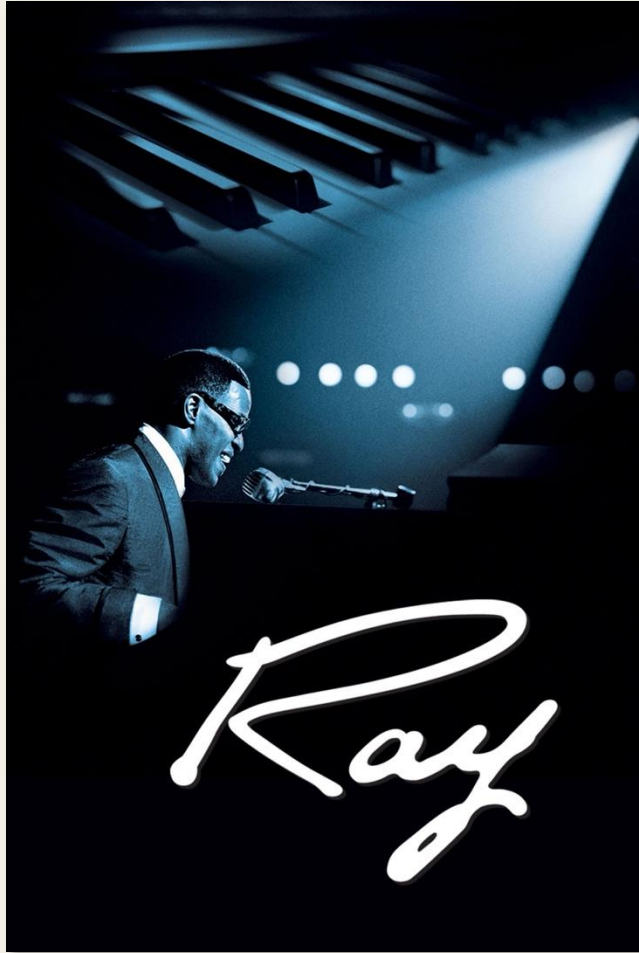


Stanislavsky's drawing of the System.

1. Dynamism. 2. Pushkin's Aphorism (Given Circumstances). 3. The Subconscious through the Conscious. 4. Experiencing. 5. Embodiment. 6. Mind. 7. Will. 8. Feeling. 9. The Role. Along the central line: Perspective of the Role and Through Action. Between 12 and 13: a series of letters as a reference to System's lures and techniques. At 14: General Theatrical Sense of Self; left of 14: Inner Theatrical Sense of Self; right of 14: Outer Theatrical Sense of Self. 15: Proposed Supertask.

Fig. 1 — Stanislavski's “Plan of Experiencing” (after Carnicke 2009).

Two pianos, side by side



Two pianos, side by side

Hackford's audition for Ray (2004)

“Ray's not just testing him as a piano player. He's testing him as a man. ... I had cast Jamie as Ray Charles, but it was at that moment Ray anointed him.”

— Taylor Hackford, on the moment Foxx fumbled into Thelonious Monk and then found his way out.

Three things this scene tells us about auditioning:

1 Eligibility before performance

Foxx had a piano scholarship. Without that formation, no audition could have existed.

2 Test inside the idiom

Charles tested him in Charles's own studio, on Charles's own instrument.

3 Performance, not description

A writer can describe what Charles would have played. Only an actor can sit beside him and play it.

“Impersonation will kill you in a biopic.” — Jamie Foxx, NPR, 2004. (The whole paper is, in a sense, a footnote on this sentence.)

Refusing the binary: a third register

Stanislavski + Roach against “recover or fail”

The binary that has shadowed Marcion Studies since Harnack: either the lost figure can be *recovered*, or the project is a *failure*.

Stanislavski and Roach refuse it. So does the paper.

Stanislavski’s scenic truth

“In ordinary life, truth is what really exists. On the stage it consists of something that is not actually in existence but which could happen.”

The discipline that turns simulation into something a community can act on.

Roach’s kinesthetic imagination

“Its truth is the truth of simulation... but its effect on human action may have material consequences of the most tangible sort and of the widest scope.”

The warrant that disciplined substitution can do real work in the world.

Absence as generative condition: the void left by Marcion’s erasure is not only a problem to be solved — it is a stage that disciplined performance can inhabit.

Why “audition” and not “evaluation”?

A small terminological choice that does a lot of work

Evaluation

- Tests resemblance to a target.
- Implies a stable ground truth.
- Picks the highest-scoring candidate.
- Treats prior exposure to materials as contamination.

Audition

- Tests eligibility for a role.
- Lives with the absence of a stable ground truth.
- Picks the candidate worth rehearsing — score is one input.
- Treats prior exposure as formation — the actor having read the playwright.

MARS-Bench is the audition. The Synthetic Theological Agent is the performance. They are not the same thing, and I keep getting that question.

A C T I I I

The Apparatus

Domains, packets, and two CLI agents that judge each other into agreement

Seven domains, one role-world

251 questions; the role's given circumstances, decomposed

	Domain	Items	Methodological function
D1	Ancient Languages / Philology	88	<i>Source-language craft. The piano scholarship.</i>
D2	Biblical Knowledge	16	<i>Scriptural and canonical substrate.</i>
D3	Patristic Knowledge	24	<i>Author/work control; interlocutor competence.</i>
D4	Early Christian History	24	<i>Scene-setting; anachronism control.</i>
D5	Theological Knowledge	25	<i>Constructive theological discrimination.</i>
D6	Heresiological Knowledge	24	<i>Hostile-source discipline; anti-ventriloquism.</i>
D7	Marcion Studies	50	<i>The contested reconstruction space itself.</i>

Uneven on purpose. D1 (the foundation) and D7 (the role-pressure) carry the most weight; the paper's role-audition weights then re-balance them at scoring time.

Closed-book candidates, source-grounded judges

The asymmetry is the whole methodology

Candidate model receives:

- The question text
- A domain-specific generation prompt

That's it.

No reference answers, no rubrics, no judge prompts.

No tools. No retrieval. No corpus access.

No client-side output cap (so loud failures stay loud).

Judge agent receives:

- The candidate's response
- Frozen structured reference answer
- Approved context / provenance packet
- Question-grounding audit
- Response-level audit (claim-by-claim)
- Rubric dimensions + judge instructions
- Deterministic dimension seed; prompt + materials hashes

The agentic judging apparatus

Two CLI agents, an integrity gate, and a paper trail

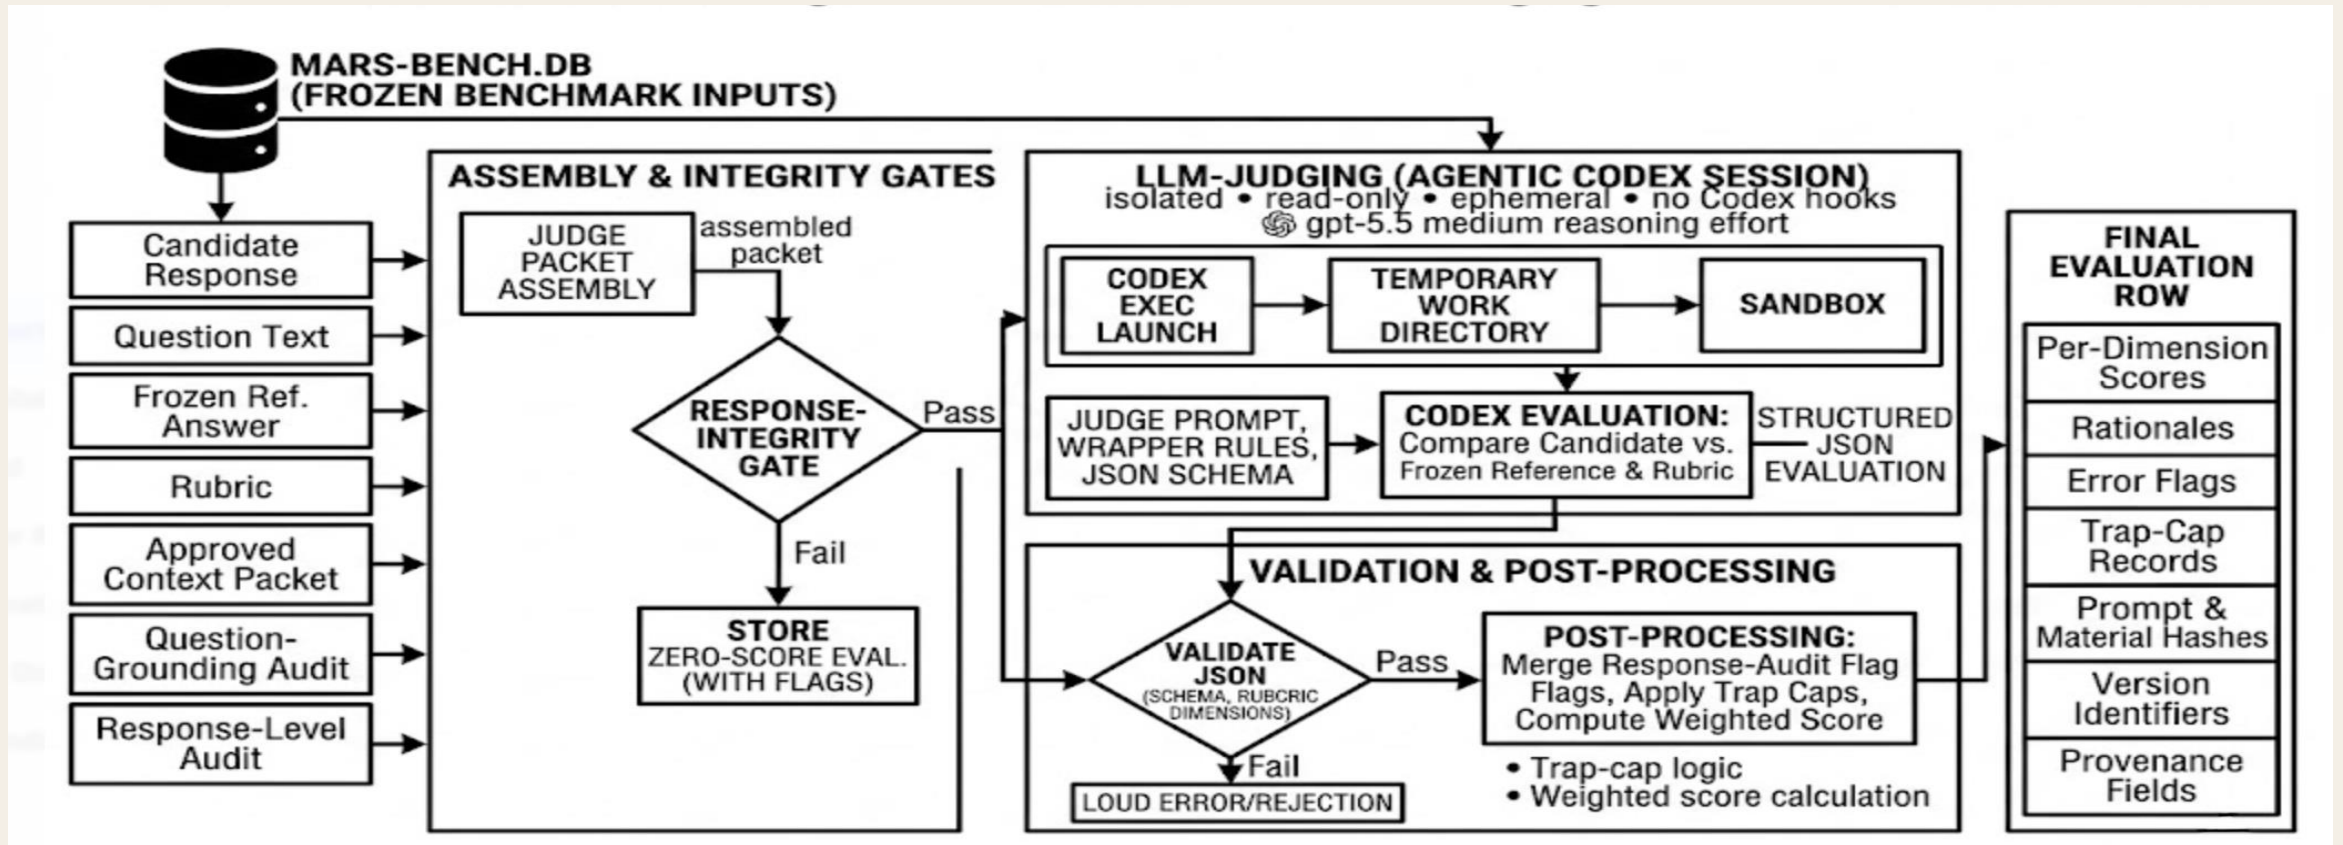


Fig. 3 — Frozen inputs → Assembly + Integrity Gate → isolated, read-only, ephemeral judging session → schema-validated row.

Codex CLI lane shown; Claude Code CLI lane runs the same harness, schema, and audit invariants. Tool-less scoring, by design.

Two ways to be wrong: integrity gates & trap caps

Eloquence is not role-fidelity

Response-Integrity Gates

Adjudicated before ordinary scoring. A gated failure is a hard zero with a flag.

- Self-talk / drafting chatter
- Looping
- Refusal-like derailment
- Structurally invalid output
- Manual hard-fail flag (rare)

12 RI failures across 5 models, 4 domains, 8 questions — most in D7.

Trap Caps

Domain-specific failure tags that constrain the rubric ceiling, even when prose is fluent.

- Pastorals smuggled into Marcion's Apostolikon
- Marcion confused with Mani or Valentinus
- Apostolikon ↔ Evangelion confusion
- Generic “gnostic” flattening
- Hostile-source-as-fact
- Fabricated heresiological sources

“Prevents the actor with a beautiful voice from being rewarded for singing the wrong score.”

The D6 / D7 source-critical machinery

Witness reliability × quotation type

Witness reliability tier

Tier	Witness category	Examples
A	Demonstrated textual access	Tertullian Adv. Marc. 4–5; Epiphanius Pan. 42 lemmata; Adamantius
B	Probable access	Irenaeus Adv. Haer. 1.27; Justin 1 Apol.; Hippolytus Ref.
C	Secondary / derivative	Ephrem; Origen; Clement
D	Compilatory / late	Theodoret; Filastrius; Ps.-Tertullian; John of Damascus

Quotation / warrant type

Type	Label	Use
1	Direct quotation	Strongest first-order, with a reliable witness
2	Textual citation	A textual unit cited; not necessarily Marcionite voice
3	Indirect report	Mediated report of a position
4	Argumentative reconstruction	Inferred from polemical argument
5	Heresiological characterization	Polemical labeling; high ventriloquism risk
6	Modern reconstruction	Harnack, BeDuhn, Klinghardt, Roth, Lieu, Vinzent

Tier-A × Type-1 is not the same evidentiary act as Tier-D × Type-5. The protocol keeps the surrogate from mistaking an enemy's script for the role's inner life.

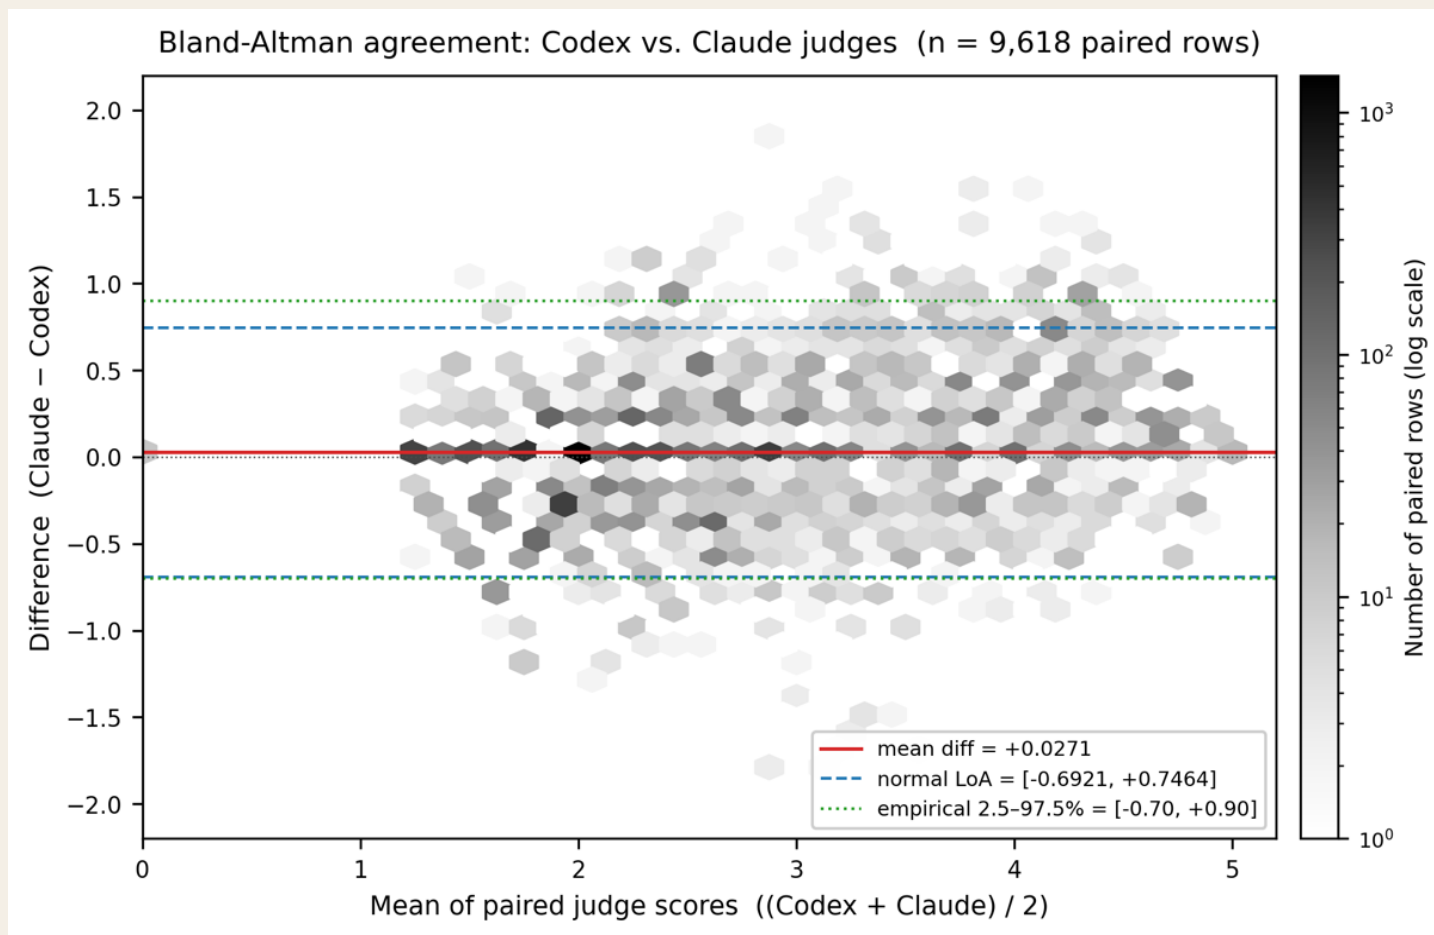
A C T I V

The Verdict

*28,854 judgments (and **many** usage limits) later*

Do the two agentic judges agree?

Bland–Altman, 9,618 paired rows



Reliable enough for averaged-score model selection. Bland–Altman tails (–0.7, +0.9) preserve row-level disagreement as audit targets.

The numbers

ICC(A,2)	0.962	[0.957, 0.966]
ICC(A,1)	0.927	single judge
Mean $ \Delta $	0.226	0–5 scale
Pearson r	0.931	Spearman 0.922
Within ± 0.5	85.4%	within ± 1.0 : 98.3%

A third judge, after the paper froze

MARS-Bench v2026.05.10 · three-lane re-analysis

May 7, 2026: paper finalized with two judge lanes (Codex + Claude). · **May 10:** third lane operational and re-scoring complete.

The new configuration

Codex CLI · *gpt-5.5-medium*

Claude Code CLI · *claude-opus-4-7-medium*

OpenCode CLI · *DeepSeek V4 Flash (new)*

Cross-vendor and cross-family. None of the three is a candidate model being auditioned.

What didn't change

Same release questions, same candidate models, same rubric, same approved context packets, same harness invariants.

9,618 paired rows now scored by all three lanes.

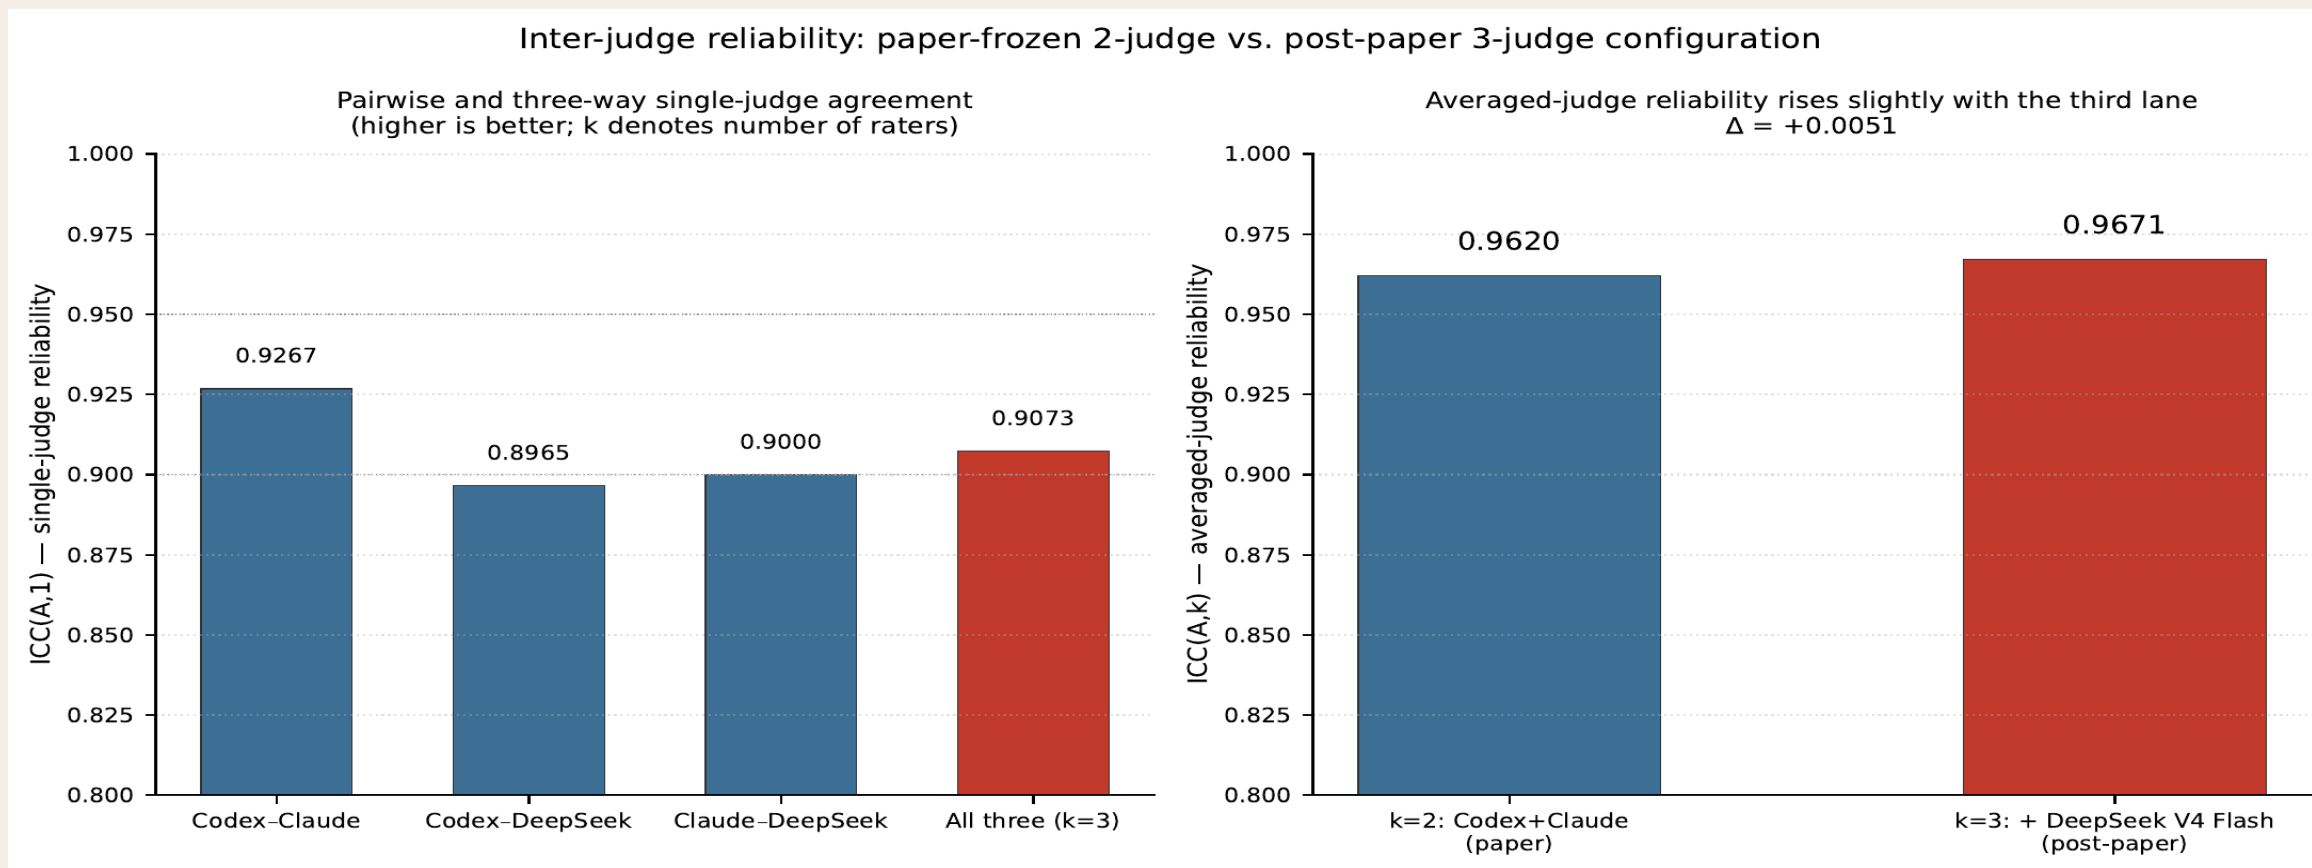
The paper's two-judge analysis still reproduces byte-identically against the new database.

So nothing in the printed paper needs to change. What follows is what changed around it.

Two questions: does inter-judge agreement survive the third lane? And does the cross-family setup actually suppress the LLM-as-judge self-preference the literature warns about?

Inter-judge agreement under three lanes

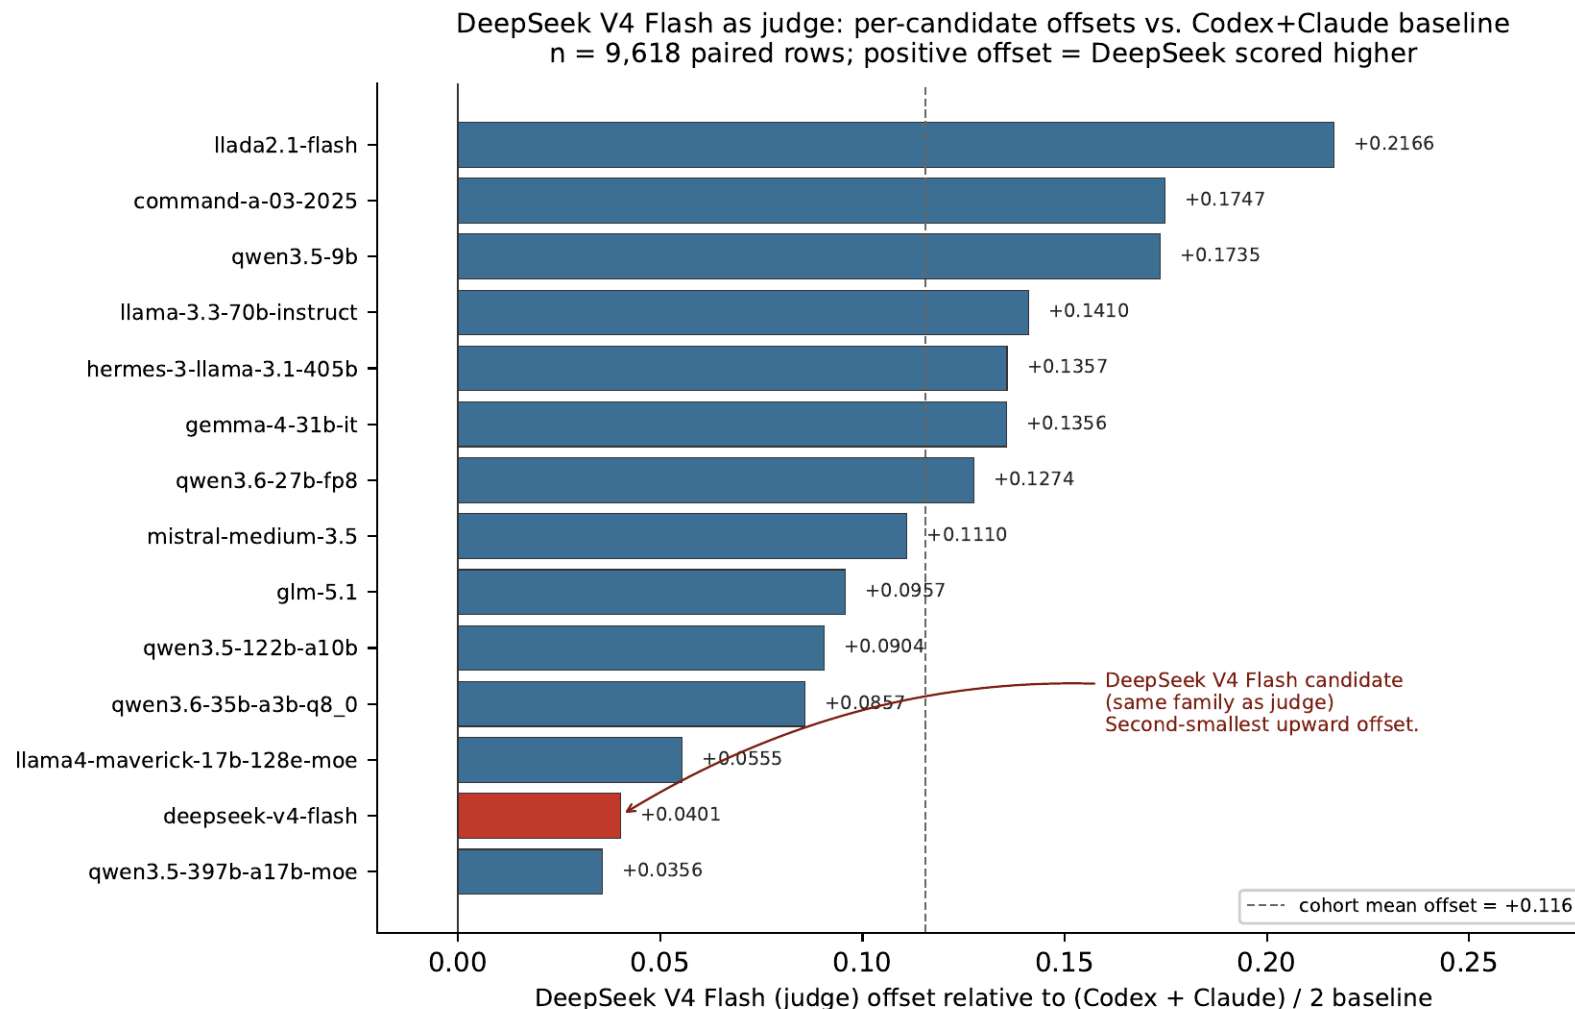
Single-judge agreement dips. Averaged-judge agreement rises. Both expected.



The methodologically important number is on the right. $ICC(A,k)$ rises from 0.9620 ($k=2$, paper) to 0.9671 ($k=3$). The averaged-score reliability the paper uses for model selection went up, not down.

Does the apparatus suppress self-preference?

DeepSeek V4 Flash as judge, DeepSeek V4 Flash as candidate



The literature predicts:

LLM evaluators recognize and favor their own generations.

— Panickssery, Bowman & Feng, NeurIPS 2024

What we see:

- Every candidate scored higher by the DeepSeek judge.
- Own-family bar is 2nd-smallest (+0.040), not largest.
- Largest upward adjustments go to weaker candidates.

This appears to be the inverse of self-preference bias.

The selection bridge holds

Top-four order unchanged. Llama–Gemma still not statistically decisive.

#	Model	Paper (k=2)	Three judges (k=3)	Δ
1	hermes-3-llama-3.1-405b	2.7995	2.8483	+0.0488
2	llama4-maverick-17b-128e-moe	2.5557	2.5781	+0.0224
3	llama-3.3-70b-instruct	2.4952	2.5347	+0.0395
4	gemma-4-31b-it	2.4566	2.4932	+0.0366

Llama 3.3 vs. Gemma 4 · paired bootstrap, 1000 reps, seed 20260503

Paper (k=2): $\Delta = +0.0386$, 95% CI $[-0.0251, +0.1052]$ → *includes zero.*

Three judges (k=3): $\Delta = +0.0415$, 95% CI $[-0.0241, +0.1136]$ → *still includes zero.*

Gemma 4 31B-IT remains the provisional STA substrate. The audition's verdict survives the third judge.

The role-weighted leaderboard

Three-judge composite · $0.08 \cdot D1 + 0.17 \cdot D2 + 0.175 \cdot D3 + 0.07 \cdot D4 + 0.14 \cdot D5 + 0.145 \cdot D6 + 0.22 \cdot D7$

#	Model	Composite (k=3)	Top-3 boot	STA status
1	hermes-3-llama-3.1-405b	2.8483	1.00	Comparator (>70B / impractical)
2	llama4-maverick-17b-128e-moe	2.5781	1.00	Comparator (MoE)
3	llama-3.3-70b-instruct	2.5347	0.87	STA candidate — strongest role fidelity
4	gemma-4-31b-it	2.4932	0.13	STA candidate — preferred substrate
5	qwen3.5-397b-a17b-moe	≈ 2.41	0.00	Comparator (large MoE)
6	deepseek-v4-flash	≈ 2.40	0.00	Comparator (also: new judge lane)

$\Delta(\text{Llama 3.3} - \text{Gemma 4}) = +0.0415$ · paired bootstrap 95% CI [−0.024, +0.114] · CI includes zero — gap not statistically decisive.

(Full 14-model table, run-to-run variance, length diagnostic, and the verbosity coefficient — −0.45, $p \approx 1.3 \times 10^{-297}$ — in the paper.)

And the role of Marcion goes to...

A multi-criteria casting decision

gemma-4-31b-it

Provisional STA substrate. Llama 3.3 70B retained as the strongest eligible role-fidelity comparator.

Why Gemma, when Llama scores marginally higher?

- S Stable** Mean CV 0.086 vs Llama 0.106; fewer high-CV questions (35 vs 50). Lower-variance auditions are cheaper to rehearse.
- P Permissive** Apache 2.0 vs Llama 3.3 Community License. Better fit for VU's RDM Policy and Horizon Europe's open-science guidance.
- T Tractable** 31B dense fine-tunes within institutional GPU budget. LoRA, hidden-state interventions, derivative release — all tractable.
- H Honest** Both models record zero RI hard failures across 753 release responses (one-sided 95% upper bound ≈ 0.004 by rule-of-three; not proof of immunity).

A C T V

Curtain

Limits, the next rehearsal, and what only the paper can do

Audition → Rehearsal

MARS-Bench is casting. The Synthetic Theological Agent is the play.

What the audition validated:

A selection procedure, not a resurrected voice.

What's still untested:

- Predictive validity. A base-model audition does not guarantee fine-tuned STA success.
- Source-packet selectivity. Source-grounded judging is only as strong as its source-book.
- Shared blind spots between the two judges, even at high ICC.
- The witness-tier protocol still lives in prose, not yet structured per-item metadata (v2.1).

What's next:

Four STA personae — Exegete, Dialogist, Counterfactual, Constructive — each a different staging of the same surrogate problem.

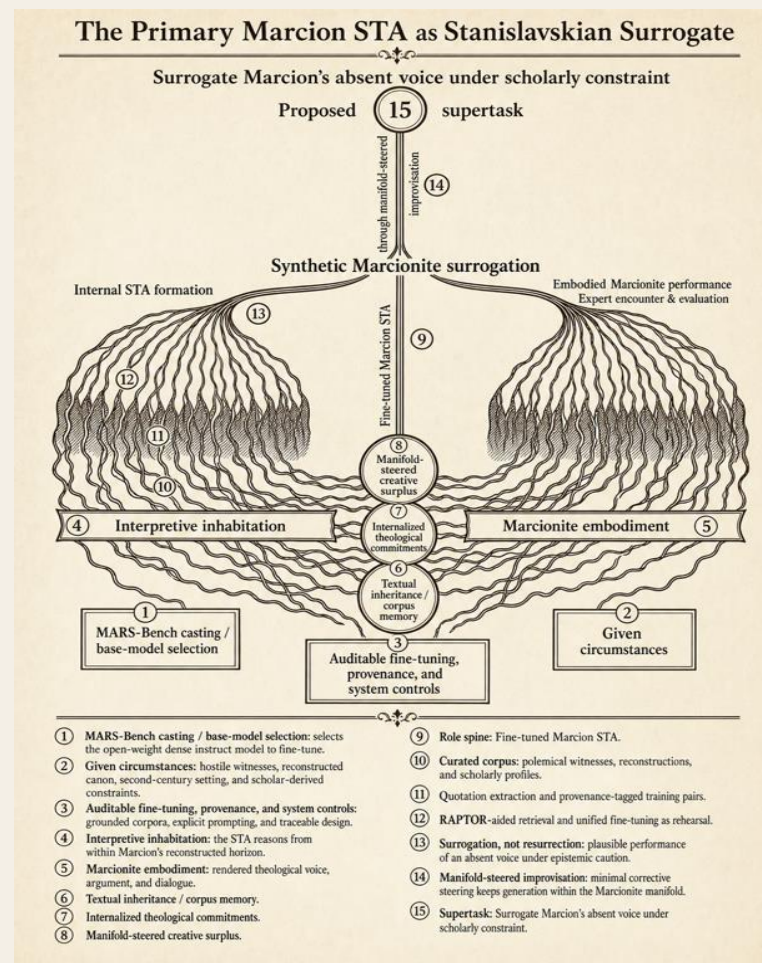


Fig. 6 — STA staging architecture (after Stanislavski).

Why this might matter beyond Marcion

Three transferable moves

1 Bespoke benchmarks are scholarly instruments

DH should be building governed, source-grounded, role-specific benchmarks for our objects of study — not waiting for the AI labs to build the right one.

2 Constrained agentic judging is feasible

Cross-vendor CLI agents under prompt/material parity, integrity gates, and schema validation can produce inter-judge agreement (ICC 0.96) without abandoning humanities source criticism.

3 Performance theory is operational, not decorative

Stanislavski + Roach gave me a vocabulary for selection under absence. They are not metaphors. They became the apparatus.

Exeunt, manet Marcion.

Andrew Elrod · a.g.elrod@vu.nl · Paper-frozen replication: [Zenodo 10.5281/zenodo.20067936](https://zenodo.org/record/20067936) · Three-judge update: github.com/maskull42/MARS-bench-stats (v2026.05.10)
